

## HW3

### 1 Homework 3: This HW is based on the code for Lecture 7 (Gradient Descent).

#### 1.1 Instructions:

Place the answer to your code only in the area specified. Also, make sure to run all your code, meaning, press » to “Restart Kernel and Run All Cells”. This should plot all outputs including your answers to homework questions. After this, go to file (top left) and select “Print”. Save your file as a PDF and upload the PDF to Canvas.

#### 2 Question:

Try to build a simple optimizer to minimize:

$$f(w) = a[0] + a[1]*w + a[2]*w^2 + \dots + a[d]*w^d$$

for the coefficients  $a = [0, 0.5, -2, 0, 1]$ .

- Plot the function  $f(w)$  (2 points)
- Can you see where the minima is? (1 point)
- Write a function that outputs  $f(w)$  and its gradient (3 points).
- Run the optimizer on the function to see if it finds the minima (3 points).
- Print the function value and number of iterations (3 points).
- Instead of writing the function for a specific coefficient vector  $a$ , create a class that works for an arbitrary vector  $a$  (3 points).

You may wish to use the `poly.polyval(w,a)` method to evaluate the polynomial.

```
[ ]: import asyncio
import typing
import numpy.polynomial.polynomial as poly
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
%matplotlib inline

a = [0, 0.5, -2, 0, 1]

print(poly.polyval(-1, a))
print(poly.polyval(0, a))
```

```
print(poly.polyval(0.25, a))
print(poly.polyval(1, a))
print(poly.polyval(2, a))
```

```
-1.5
0.0
0.00390625
-0.5
9.0
```

Checking by hand

$f(w) = a[0] + a[1]*w + a[2]*w^2 + \dots + a[d]*w^d$  for the coefficients  $a = [0, 0.5, -2, 0, 1]$ .

$$f(w) = \frac{1}{2}w - 2w^2 + 0 + w^4$$


$$w = 1; 0.5 - 2 + 1 = 0.5 - 1 = -0.5$$

$$w = 2; 1 - 8 + 16 = 1 + 8 = 9$$

## 2.1 Plot the function $f(w)$ (2 points)


```
[ ]: w = np.linspace(-10, 10, 100)
      v = poly.polyval(w, a)

      plt.plot(w,v)
      plt.grid(True)
```



2.2 Can you see where the minima is? (1 point)

```
[ ]: plt.plot(w,v)
      plt.xlim(-2.5, 2.5)
      plt.ylim(-2, 2)
      plt.grid(True)
```



Looking at the macro graph, one could assume  $f(0)$  would have been the minima but this is not true. Zooming in to the graph  $f(\approx -1)$  is the true minima while  $f(\approx 1)$  is a local minima.

### 2.3 Write a function that outputs $f(w)$ and its gradient (3 points).

```
[ ]: p = np.polynomial.Polynomial(a)
print(p)

d = p.deriv()
print(d)

assert p(1) == -0.5, "Our polynomial is wrong..."
```

```
0.0 + 0.5 x - 2.0 x**2 + 0.0 x**3 + 1.0 x**4
0.5 - 4.0 x + 0.0 x**2 + 4.0 x**3
```

```
[ ]: # our gradient function
feval = lambda w: (p(w), p.deriv()(w))

w0 = np.random.randn(1)

# Perturb the point
```

```

step = 1e-6
w1 = w0 + step*np.random.randn(1)

# Measure the function and gradient at w0 and w1
f0, fgrad0 = feval(w0)
print(f0, fgrad0)
f1, fgrad1 = feval(w1)
print(f1, fgrad1)

# Predict the amount the function should have changed based on the gradient
df_est = fgrad0.dot(w1-w0)
print(df_est)

# Print the two values to see if they are close
print("Actual f1-f0      = %12.4e" % (f1-f0))
print("Predicted f1-f0 = %12.4e" % df_est)

assert np.isclose(f1-f0, df_est), "Actual and Predicted value do not match!"

```

```

[-0.39107492] [1.53090923]
[-0.39107624] [1.53089991]
-1.3229735478990994e-06
Actual f1-f0    = -1.3230e-06
Predicted f1-f0 = -1.3230e-06

```

## 2.4 Run the optimizer on the function to see if it finds the minima (3 points).

```

[ ]: async def grad_opt_sim(feval: typing.Callable, winit: np.array, lr: float=1e-3,
    ↪ nit: int=1000) -> typing.Tuple:
    """
    Simple gradient descent optimization

    feval: A function that returns f, fgrad, the objective
           function and its gradient
    winit: Initial estimate
    lr:    learning rate
    nit:   Number of iterations
    """
    # Initialize
    w0 = winit
    f0, fgrad0 = feval(w0)

    # Loop over iterations
    for it in range(nit):

        # Evaluate the function and gradient
        temp_f0, fgrad0 = feval(w0)

```

```

# https://numpy.org/doc/stable/reference/generated/numpy.isclose.html
# using to determine if we aren't making "enough" progress
if np.isclose(temp_f0, f0, rtol=1e-05, atol=1e-08) and it > 0:
    break
f0 = temp_f0

# Take a gradient step
w0 = w0 - lr*fgrad0

return (w0, f0, it)

# We need multiple random points to find a true minima vs local minma
winits = np.array([np.random.randn(1)[0] for i in range(15)])

data = asyncio.gather(*[grad_opt_sim(feval, winit) for winit in winits])
await data

```

```

[ ]: [(0.6367990229171355, -0.3271556089271581, 999),
      (-0.9937295213879747, -1.496402714509664, 999),
      (-1.070086776260489, -1.5139786632587726, 356),
      (0.9134230000018356, -0.5158346054591425, 999),
      (0.9413653484936304, -0.5163546200852214, 387),
      (-1.0443045777344577, -1.513933986900044, 874),
      (0.7405611127477955, -0.42510027589677113, 999),
      (-1.044276005002227, -1.5139304429142142, 544),
      (-1.0701016676897654, -1.5139768244001859, 305),
      (-1.0443148721131452, -1.513935261901335, 464),
      (0.5032053741302854, -0.18970560573021664, 999),
      (0.9413406378715726, -0.5163564053353972, 109),
      (-1.040450163973884, -1.513387974660023, 999),
      (0.8958712125290627, -0.5130484076585351, 999),
      (-1.0442434880471358, -1.513926400526765, 436)]

```

## 2.5 Print the function value and number of iterations (3 points).

```

[ ]: df = pd.DataFrame(columns=["init_x", "init_y", "grad_x", "grad_y"])
df["init_x"] = winits
df["init_y"] = p(winits)
df[["grad_x", "grad_y", "iters"]] = data.result()

df

```

```

[ ]:
      init_x  init_y  grad_x  grad_y  iters
0   0.144875  0.030900  0.636799 -0.327156  999.0
1   0.077826  0.026836 -0.993730 -1.496403  999.0
2  -1.837919  3.735669 -1.070087 -1.513979  356.0

```

```

3  0.259492 -0.000392  0.913423 -0.515835  999.0
4  1.097689 -0.409162  0.941365 -0.516355  387.0
5 -0.031867 -0.017963 -1.044305 -1.513934  874.0
6  0.154071  0.030123  0.740561 -0.425100  999.0
7 -0.433276 -0.556853 -1.044276 -1.513930  544.0
8 -1.383997 -0.853956 -1.070102 -1.513977  305.0
9 -0.593715 -0.877598 -1.044315 -1.513935  464.0
10 0.137966  0.031276  0.503205 -0.189706  999.0
11 0.952835 -0.515099  0.941341 -0.516356  109.0
12 0.040085  0.016832 -1.040450 -1.513388  999.0
13 0.214484  0.017352  0.895871 -0.513048  999.0
14 -0.649219 -0.989930 -1.044243 -1.513926  436.0

```

## 2.6 Instead of writing the function for a specific coefficient vector $\mathbf{a}$ , create a class that works for an arbitrary vector $\mathbf{a}$ (3 points).

```

[ ]: # It should already work for any arbitrary coef vector

class GradientPolynomial(object):
    def __init__(self, a: np.array) -> None:
        self.__p = np.polynomial.Polynomial(a)
        self.__inits = []
        self.__vals = []

    def __repr__(self) -> str:
        return f"<GradientPolynomial: f{self.__p}>"

    @property
    def __d(self) -> np.polynomial.Polynomial:
        return self.__p.deriv()

    def feval(self, w: np.array) -> typing.Tuple:
        return (self.__p(w), self.__d(w))

    async def optimizer(self, winit: np.array, eta=1e-3, nit=1000) -> typing.
↳ Tuple:
        w0 = winit
        f0, fgrad0 = self.feval(w0)

        for it in range(nit):
            temp_f0, fgrad0 = self.feval(w0)

            # https://numpy.org/doc/stable/reference/generated/numpy.isclose.
↳ html
            # using to determine if we aren't making "enough" progress
            if np.isclose(temp_f0, f0, rtol=1e-07, atol=1e-08) and it > 0:
                break

```

```

        f0 = temp_f0

        w0 -= eta*fgrad0

    return (w0, f0, it)

async def plot(self,
                n: int=1000,
                domain: tuple=(-10, 10),
                initial_points: typing.Union[np.array, int, None]=None,
                eta: float=1e-3,
                nit: int=1000,
                xlim: typing.Union[tuple, None]=None,
                ylim: typing.Union[tuple, None]=None) -> pd.DataFrame:
    x, y = self.__p.linspace(n, domain=domain)

    if isinstance(initial_points, int):
        self.__inits = np.random.choice(x, initial_points)
        self.__vals = []
    elif isinstance(initial_points, (np.array, list, tuple)):
        self.__inits = initial_points
        self.__vals = []
    elif not self.__inits:
        raise Exception("No array of initial points or number of initial
↳points to be generated was provided")

    if not self.__vals:
        self.__vals = asyncio.gather(*[self.optimizer(i, eta, nit) for i in
↳self.__inits])
        await self.__vals

    df = pd.DataFrame(columns=["init_x", "init_y", "grad_x", "grad_y"])
    df["init_x"] = self.__inits
    df["init_y"] = self.__p(self.__inits)
    df[["grad_x", "grad_y", "iters"]] = self.__vals.result()

    plt.plot(x, y)
    plt.plot(*self.__d.linspace(n, domain=domain))
    plt.scatter(df["init_x"], df["init_y"], marker='o', c='black', s=30)
    plt.scatter(df["grad_x"], df["grad_y"], marker='o', c='red', s=30)

    if xlim:
        plt.xlim(*xlim)
    if ylim:
        plt.ylim(*ylim)
    plt.grid(True)

```



```

    return df

gp = GradientPolynomial(a)
print(gp)
await gp.plot(domain=(-2, 2), initial_points=20, xlim=(-3, 3), ylim=(-3, 3))


```

<GradientPolynomial:  $f0.0 + 0.5 x - 2.0 x^2 + 0.0 x^3 + 1.0 x^4$ >

```

[ ]:      init_x    init_y    grad_x    grad_y  iters
0  -1.871872  4.333601 -1.058786 -1.514745  594.0
1  -1.179179 -1.437128 -1.058784 -1.514745  461.0
2   0.410410 -0.103297  0.926385 -0.516697  999.0
3   1.291291  0.091113  0.931616 -0.516744  809.0
4   1.203203 -0.197965  0.931616 -0.516744  782.0
5   1.735736  3.919145  0.931620 -0.516744  867.0
6   0.582583 -0.272320  0.929096 -0.516743  999.0
7  -1.947948  5.835263 -1.058783 -1.514745  598.0
8  -1.239239 -1.332630 -1.058783 -1.514745  496.0
9   1.471471  1.093493  0.931617 -0.516744  842.0
10 -1.307307 -1.150898 -1.058791 -1.514745  521.0
11 -1.083083 -1.511589 -1.058786 -1.514745  309.0
12 -0.990991 -1.495174 -1.056123 -1.514745  423.0
13 -0.214214 -0.196777 -1.056123 -1.514745  923.0
14  0.782783 -0.458645  0.929180 -0.516744  793.0
15 -0.242242 -0.235040 -1.056118 -1.514745  902.0
16 -0.530531 -0.748969 -1.056114 -1.514745  738.0
17  0.166166  0.028623  0.812140 -0.477669  999.0
18 -0.038038 -0.021911 -1.053764 -1.514689  999.0
19 -1.023023 -1.509342 -1.056115 -1.514745  348.0


```



```
[ ]: gp = GradientPolynomial([0, 3, 2])
      print(gp)
      await gp.plot(domain=(-3, 3), initial_points=10, xlim=(-2, 1), ylim=(-2, 2))
```

<GradientPolynomial:  $f0.0 + 3.0 x + 2.0 x^2$ >


```
[ ]:      init_x    init_y    grad_x    grad_y    iters
0 -1.342342 -0.423261 -0.760762 -1.124766 999.0
1  0.237237  0.824275 -0.732063 -1.124351 999.0
2  0.795796  3.653969 -0.721914 -1.123410 999.0
3 -2.093093  2.482798 -0.774403 -1.123799 999.0
4  0.525526  2.128931 -0.726825 -1.123917 999.0
5 -0.987988 -1.011723 -0.754324 -1.124962 999.0
6 -2.585586  5.613749 -0.783351 -1.122757 999.0
7 -1.900901  1.524146 -0.770911 -1.124118 999.0
8  1.798799 11.867751 -0.703690 -1.120676 999.0
9 -2.393393  4.276484 -0.779859 -1.123202 999.0
```



```
[ ]: gp = GradientPolynomial([0, 0, -2, -3, 4, 0.5])
      print(gp)
      await gp.plot(domain=(-1, 1), initial_points=10, ylim=(-2, 2))
```

<GradientPolynomial:  $f0.0 + 0.0 x - 2.0 x^{**2} - 3.0 x^{**3} + 4.0 x^{**4} + 0.5 x^{**5}$ >

```
[ ]:      init_x  init_y  grad_x  grad_y  iters
0 -0.089089 -0.013503 -0.291243 -0.067801 999.0
1  0.431431 -0.467121  0.784172 -1.015678 460.0
2  0.451451 -0.508118  0.784170 -1.015678 450.0
3  0.417417 -0.438893  0.784172 -1.015678 467.0
4  0.653654 -0.902494  0.784173 -1.015678 346.0
5  0.517518 -0.645980  0.784177 -1.015678 419.0
6 -0.743744  1.238042 -0.296714 -0.067857 995.0
7 -0.219219 -0.055524 -0.295170 -0.067857 940.0
8  0.941942 -0.762095  0.785444 -1.015678 316.0
9  0.099099 -0.022170  0.784177 -1.015679 754.0
```



## 2.7 Extra/Alt Implementation

```
[ ]: x, y = p.linspace(1000, domain=(-3, 3))

# choose 10 random points from domain -3 to 3
initial_points = np.random.choice(x, 10)
print(f"Random initial points: {initial_points}")

async def gradient_descent(pt, eta=0.01, steps=5000, allow_early_out=True) -> float:
    result = pt
    for i in range(steps):
        temp_result = result - eta * d(result)
        if allow_early_out and (p(temp_result) >= p(result)):
            break
        result = temp_result
    return result

data = asyncio.gather(*[gradient_descent(init_pt) for init_pt in initial_points])
await data
```

```

data = np.array(data.result())

print(f"Gradient Descent x-values: {data}")
print(f"Gradient Descent y-valyes: {[p(d) for d in data]}")

# this also works for plotting
plt.plot(x, y)
plt.plot(*d.linspace(1000, domain=(-10, 10)))
plt.scatter(initial_points, p(initial_points), marker='o', c='black', s=30)
plt.scatter(data, p(data), marker='o', c='red', s=30)
plt.xlim(-3.5, 3.5)
plt.ylim(-3.5, 3.5)
plt.grid(True)

```

Random initial points: [ 1.67867868 1.1981982 -2.89189189 -1.996997  
1.88288288 -2.58558559  
0.8018018 2.24324324 0.36936937 -1.04804805]  
Gradient Descent x-values: [ 0.93040295 0.93040294 -1.05745378 -1.05745379  
0.93040296 -1.05745379  
0.93040291 0.93040295 0.93040291 -1.05745376]  
Gradient Descent y-valyes: [-0.5167485082852478, -0.5167485082852483,  
-1.5147536412757052, -1.5147536412757043, -0.5167485082852467,  
-1.5147536412757039, -0.5167485082852478, -0.5167485082852481,  
-0.5167485082852488, -1.514753641275705]

