

LECTURE 5: LASSO REGULARIZATION AND FEATURE SELECTION

Ehsan Aryafar

earyafar@pdx.edu


<http://web.cecs.pdx.edu/~aryafare/ML.html>

Recall: Bias and Variance

- To understand potential problem of using a large model class introduce two key quantities:
- **Bias:** $Bias(\mathbf{x}_{test}) := f_0(\mathbf{x}_{test}) - E[f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}})]$
 - How much the average value of the estimate differs from the true function
- **Variance:** $Var(\mathbf{x}_{test}) := E \left[f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}}) - E[f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}})] \right]^2$
 - How much the estimate varies around its average
- Bias and variance are (conceptually) measured as follows:
 - Get many independent training data sets, each with same size N and input values x_i
 - Each dataset has different output values y_i because of independent noise in the training data
 - Obtain $\hat{\boldsymbol{\beta}}$ for each training data set
 - Bias and variances are computed over the different sets
- Of course, in reality, we have only one training dataset
 - Used to study theoretical averages over different experiments

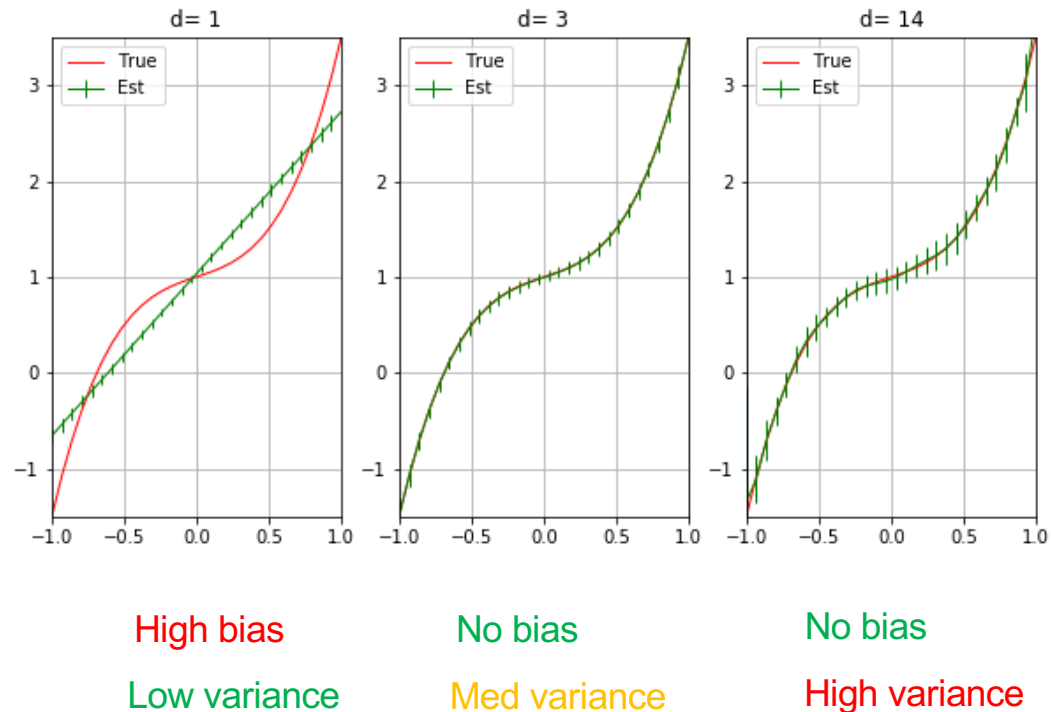
Recall: Bias Illustrated

- Red: True function
- Repeat 100 trials
 - Each trial has independent data
 - Obtain estimate for each trial
- Dashed line: Mean estimate among all trials
- Bias = True – Mean estimate
- Conclusions:
 - Low model orders \Rightarrow bias high
 - High model orders \Rightarrow bias low



Recall: Variance Illustrated

- Red: True function
- Repeat 100 trials
 - Each trial has independent data
 - Obtain estimate for each trial
- Variance=STD around mean
- Conclusions:
 - Low model orders \Rightarrow low variance
 - High model orders \Rightarrow high variance



Recall: Bias-Variance Formula

- Recall definitions:

- Function MSE: $MSE_f(\mathbf{x}_{test}) := E[f_0(\mathbf{x}_{test}) - f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}})]^2$:

- Bias: $Bias(\mathbf{x}_{test}) := f_0(\mathbf{x}_{test}) - E[f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}})]$

- Variance: $Var(\mathbf{x}_{test}) := E[f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}}) - E[f(\mathbf{x}_{test}, \hat{\boldsymbol{\beta}})]]^2$

- Bias-Variance formula : $MSE_f(\mathbf{x}_{test}) = Bias(\mathbf{x}_{test})^2 + Var(\mathbf{x}_{test})$

- Will be proved soon

- Bias-Variance tradeoff


- Bias due to under-modeling

- Reduced with high model order

- Variance is due to noise in training data and number of parameters to estimate

- Increases with higher model order

Recall: Bias-Variance Tradeoff



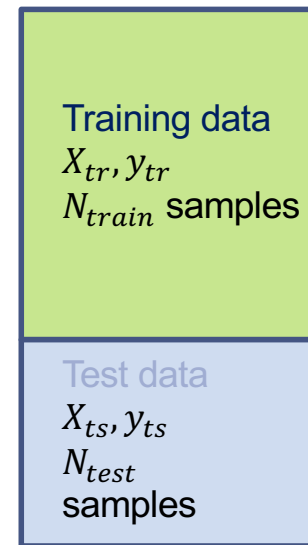
- Bias:
 - Due to under-modeling
 - Reduced with high model order
- Variance:
 - Increases with noise in training data
 - Increase with high model order
- Optimal model order depends on:
 - Amount of samples available
 - Underlying complexity of the relation

Recall: Cross Validation

- **Key idea:** Evaluate on samples different from training
- Get data X, y
- Split into training X_{tr}, y_{tr} and test X_{ts}, y_{ts}
- For $p = 1$ to p_{max} // Loop over model order
 - Fit on training data with model order p
 - Predict values on test data
 - Score fit on test data (e.g. measure RSS)
- Select model order with smallest score:

$$\hat{p} = \arg \min_p S[p]$$

- Maximize if higher score is better



$$\hat{\beta} = \text{fit}(X_{tr}, y_{tr}, p)$$

$$\hat{y}_{ts} = \text{predict}(X_{ts}, \hat{\beta})$$

$$S[p] = \text{score}(y_{ts}, \hat{y}_{ts})$$

Recall: K-Fold Cross Validation

- *K*-fold cross validation
 - Divide data into *K* parts
 - Use *K* – 1 parts for training. Use remaining for test.
 - Average over the *K* test choices
 - More accurate, but requires *K* fits of parameters
 - Typical choice: *K*=5 or 10
 - Average MSE over *K* folds estimates the total MSE
 - ($= \text{Bias}^2 + \text{Variance} + \text{irreducible error}$)
- Leave one out cross validation (LOOCV)
 - Take $K = N$ so one sample is left out.
 - Most accurate, but requires *N* model fittings
 - Necessary when *N* is small



From
<http://blog.goldenhelix.com/goldenadmin/cross-validation-for-genomic-prediction-in-svs/>

Learning Objectives

- Describe **model selection** and identify when it may be needed
- Mathematically describe linear regression with **regularization**
- Select regularizers to impose constraints such as sparsity
- Compute an L1-regularized estimate (LASSO) using sklearn
- Compute the optimal regularization level using cross validation
- Interpret results from a LASSO path
- Set regularizer based on a probabilistic prior

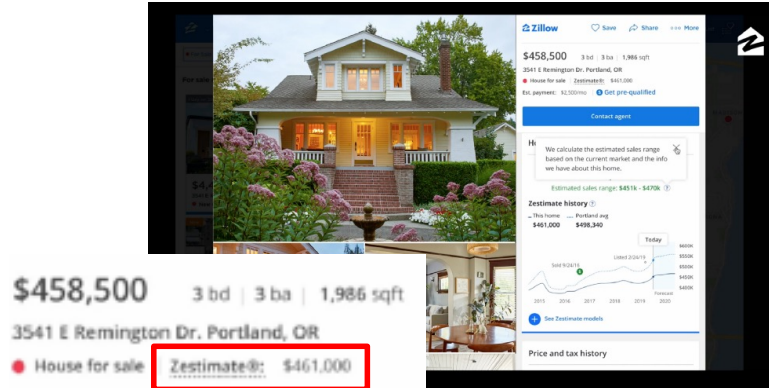
Outline

- Motivating Example: Feature selection in predicting house prices
- Model selection and regularization
- Housing price prediction with LASSO

Predicting Housing Prices

AI, MACHINE LEARNING & RESEARCH

Introducing a new and improved Zestimate algorithm



<https://www.zillow.com/tech/introducing-a-new-and-improved-zestimate-algorithm/>

- Many services now predict house prices
- Data science enters real estate!
- Many possible variables:
 - Square meters
 - Condition
 - Zip code
 - Education quality
 - ...
- What variables *really* determine the price?



Ames, Iowa Dataset



Ames, Iowa: Alternative to the Boston Housing Data as an End of Semester Regression Project

[Dean De Cock](#)

Truman State University

Journal of Statistics Education Volume 19, Number 3(2011),
www.amstat.org/publications/jse/v19n3/decock.pdf

Copyright © 2011 by Dean De Cock all rights reserved. This text may be freely shared among individuals, but it may not be republished in any medium without express written consent from the author and advance notification of the editor.

Key Words: Multiple Regression; Linear Models; Assessed Value; Group Project.

Abstract

This paper presents a data set describing the sale of individual residential property in Ames, Iowa from 2006 to 2010. The data set contains 2930 observations and a large number of explanatory variables (23 nominal, 23 ordinal, 14 discrete, and 20 continuous) involved in assessing home values. I will discuss my previous use of the Boston Housing Data Set and I will suggest methods for incorporating this new data set as a final project in an undergraduate regression course.

- Ames, Iowa Dataset
 - Sales from 2006 to 2010
 - From Dean De Cock
 - Undergraduate student
- Alternative to Boston Housing dataset
- Many more variables to explore
 - Approximately 81 variables
 - 2930 samples

Loading the Dataset

```
1 df = pd.read_csv('housing_train.csv')
2 df.head()
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscVal
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN	NaN	0
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN	NaN	0
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0

5 rows × 81 columns

- Issues:
 - Many different types of data: Discrete and continuous
 - Missing values (NaN)

Data Cleaning

```
1 nsamp, natt = df.shape
2 print('Number samples = %d' % nsamp)
3 print('Number attributes per sample = %d' % natt)
```

Number samples = 1460

Number attributes per sample = 81

```
df = df.dropna(axis=1)

nsamp, natt = df.shape
print('Number samples = %d' % nsamp)
print('Number attributes per sample = %d' % natt)
```

Number samples = 1460

Number attributes per sample = 62

```
1 df = df.loc[df['SaleCondition'] == 'Normal']
2
3 nsamp, natt = df.shape
4 print('Number samples = %d' % nsamp)
5 print('Number attributes per sample = %d' % natt)
```

Number samples = 1198

Number attributes per sample = 62

- Original data
- Remove columns with NaN values
 - Could use more sophisticated methods
- Keep only normal sales
 - Recommended in De Cock paper
 - Makes fitting much easier

Categorical Variables

- Data has many categorical variables
- Need to code the categorical variables to numerical values

Real valued

Categorical

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	...	PoolArea
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPub	...	0
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPub	...	0
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPub	...	0
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPub	...	0
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPub	...	0

Splitting the Variables

- First, split the variables into categorical and real
- Attributes are of different types (int64, object) and you can check their types

```
# Remove the ID, month sold and sales price (it is the target)
ignore_vars = ['Id', 'MoSold', 'SalePrice']

# Find real and categorical variables
cols = df.columns
cat_vars = []
real_vars = []

for col in cols:
    if not (col in ignore_vars):
        if df.dtypes[col] == 'object':
            cat_vars.append(col)
        else:
            real_vars.append(col)
```

Categorical variables have type *object*

Categorical variables = ['MSZoning', 'Street', 'Alley', 'LotShape', 'LandContour', 'Utilities', 'LotConfig', 'LandSlope', 'Neighborhood', 'Condition1', 'Condition2', 'BldgType', 'HouseStyle', 'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType', 'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinType2', 'Heating', 'HeatingQC', 'CentralAir', 'Electrical', 'KitchenQual', 'Functional', 'FireplaceQu', 'GarageType', 'GarageFinish', 'GarageQual', 'GarageCond', 'PavedDrive', 'PoolQC', 'Fence', 'MiscFeature', 'SaleType', 'SaleCondition']

Real variables = ['MSSubClass', 'LotFrontage', 'LotArea', 'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd', 'MasVnrArea', 'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'TotRmsAbvGrd', 'Fireplaces', 'GarageYrBlt', 'GarageCars', 'GarageArea', 'WoodDeckSF', 'OpenPorchSF', 'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'MiscVal', 'YrSold']

One Hot Coding

Original

	MSZoning	Street	LotShape	LandContour	Utilities	LotConfig
0	RL	Pave	Reg	Lvl	AllPub	Inside
1	RL	Pave	Reg	Lvl	AllPub	FR2
2	RL	Pave	IR1	Lvl	AllPub	Inside
4	RL	Pave	IR1	Lvl	AllPub	FR2
5	RL	Pave	IR1	Lvl	AllPub	Inside

One Hot coded

	MSZoning_FV	MSZoning_RH	MSZoning_RL	MSZoning_RM	Street_Pave
0	0	0	1	0	1
1	0	0	1	0	1
2	0	0	1	0	1
4	0	0	1	0	1
5	0	0	1	0	1

- Use pandas `get_dummies`
- Replaces categorical variables with one-hot coded values
 - Ex: MSZoning
 - Becomes MSZoning_FV, MSZoning_RH, ...

```
# Get the dataframes with real and categorical variables
df_real = df[real_vars]
df_cat = df[cat_vars]

# One-hot encode the categorical variables
df_cat_enc = pd.get_dummies(df_cat, drop_first=True)
```

Scaling Data

- Split data into training and test
- Scale data
 - Remove mean for each column and divide by standard deviation
- Needed to compare coefficients
 - Ensures that all variables have same range
 - Methods need same range for all features
- Note: The scaling transform is
 - Fit on the training data
 - Transform on test data
 - Performed on training and test

```
from sklearn.model_selection import train_test_split  
  
Xtr, Xts, ytr, yts = train_test_split(X,y,test_size=0.3)
```

```
from sklearn.preprocessing import StandardScaler  
  
# Create the scaler objects  
xscal = StandardScaler()  
yscal = StandardScaler()  
  
# Fit and transform the training data  
Xtr1 = xscal.fit_transform(Xtr)  
ytr1 = yscal.fit_transform(ytr[:,None])  
  
# Transform the test data  
Xts1 = xscal.transform(Xts)  
yts1 = yscal.transform(yts[:,None])
```

Why Fit_Transform() on Training and Transform() on Test Data

- Both are methods of class `sklearn.preprocessing.StandardScaler()`
- We want scaling be applied to our test data set too but we do not want to be biased with our model
- Using the transform method we can use the same mean and variance as it is calculated from our training data to transform our test data

First Try: Linear Regression

```

from sklearn.linear_model import LinearRegression, Ridge
from sklearn.metrics import r2_score

# Fit
reg_ls = LinearRegression()
reg_ls.fit(Xtr1, ytr1)

# Training R^2
yhat1_tr = reg_ls.predict(Xtr1)
rsq_tr = r2_score(yhat1_tr, ytr1)
print('Training R^2 = %12.4e' % rsq_tr)

# Test R^2
yhat1_ts = reg_ls.predict(Xts1)
rsq_ts = r2_score(yts1, yhat1_ts)
print('Test R^2      = %12.4e' % rsq_ts)

```

```

Training R^2 = 9.3726e-01
Test R^2      = -1.0430e+20

```

- Simple idea:
 - Use linear regression over features
- Fits the training data very well!
 - $R^2 \approx 0.937$
- But, completely fails on the test data
 - $R^2 > 10^{20}$

Conditioning

- What went wrong?
- Recall LS solution is: $\hat{\beta} = (A^T A)^{-1} A^T y$
- Matrix $A^T A$ may be ill-conditioned
 - Eigenvalues close to zero
 - Inverse blows up
- With ill-conditioned data:
 - Training error is fine
 - But the test error blows up
- Overfits data

```
from sklearn.linear_model import LinearRegression, Ridge
from sklearn.metrics import r2_score

# Fit
reg_ls = LinearRegression()
reg_ls.fit(Xtr1, ytr1)

# Training R^2
yhat1_tr = reg_ls.predict(Xtr1)
rsq_tr = r2_score(yhat1_tr, ytr1)
print('Training R^2 = %12.4e' % rsq_tr)

# Test R^2
yhat1_ts = reg_ls.predict(Xts1)
rsq_ts = r2_score(yts1, yhat1_ts)
print('Test R^2      = %12.4e' % rsq_ts)
```

```
Training R^2 = 9.3726e-01
Test R^2      = -1.0430e+20
```

Improving Conditioning via Ridge Regression

- Standard LS solution: $\hat{\beta} = (A^T A)^{-1} A^T y$
- **Ridge Regression**: Add a conditioning term:

$$\hat{\beta} = (A^T A + cI)^{-1} A^T y$$

- c is a small positive value.
 - Makes inverse well-behaved
 - We will see this technique more later
-
- Get good test R^2

```
reg_ls = Ridge(alpha=1e-5)
reg_ls.fit(Xtr1, ytr1)
yhat1 = reg_ls.predict(Xts1)
rsq = r2_score(yts1, yhat1)
print('Test R^2      = %12.4e' % rsq)
```

Test R^2 = 0.904567

What Components Matter?

- Simple idea: Look at large coefficients
- We see variables that we may expect:
 - Square footage
 - Quality
 - Zoning
- But there are some issues
 - Some variables seem highly correlated
 - Ex: GrLivArea and 2ndFlrSF
 - KitchenQual and OverallQual
 - Do we need both?

```

1 coeff_ls = reg_ls.coef_.ravel()
2 nprint = 10
3 I = np.argsort(np.abs(coeff_ls))
4 I = np.flipud(I)
5 for i in range(nprint):
6     j = I[i]
7     print('%20s %f' % (xnames[j], coeff_ls[j]))

```


```

GrLivArea 0.277146
2ndFlrSF 0.248113
KitchenQual_TA -0.235208
KitchenQual_Gd -0.229472
ExterQual_Gd -0.190125
ExterQual_TA -0.184943
YearBuilt 0.161462
MSZoning_RL 0.159591
RoofStyle_Gable -0.141669
OverallQual 0.136913

```

What Components Do *Not* Matter?

- All coefficients are far from zero
 - Very few coefficients that can be removed
- Does this mean all variables matter?
- Model or feature selection problem:
 - *How do we find the variables that matter?*



Outline

- Motivating Example: Feature selection in predicting house prices
- Model selection and regularization
- Housing price prediction with LASSO

Model Selection via Sparsity

	MSZoning_FV	MSZoning_RH	MSZoning_RL	MSZoning_RM	Street_Pave	LotShape_IR2	LotShape_IR3	LotShape_Reg	LandContour_HLS
0	0	0	1	0	1	0	0	1	0
1	0	0	1	0	1	0	0	1	0
2	0	0	1	0	1	0	0	0	0
4	0	0	1	0	1	0	0	0	0
5	0	0	1	0	1	0	0	0	0

174 variables
after one-hot
coding

- **Model selection problem:** Need to identify the parameters that *really* matter. Two reasons to do that:
 - Help interpret results
 - Improves generalization error (less parameters, close to model order)
- **Idea:** Fit model under **sparsity constraint**:
 - Linear model: $\hat{y} = \beta_0 + \beta_1 x_1 + \dots + \beta_p x_p$
 - Feature x_j is ignored if $\beta_j = 0$
 - Try to force most $\beta_j = 0 \Rightarrow$ Model only uses a few of the variables
 - Sparse coefficient constraint

Regularized LS Estimation

- **Regularization:** General method for finding constrained solutions
 - E.g. solutions that are sparse
- Standard least squares estimation (from Lecture 3, unregularized form):

$$\hat{\beta} = \arg \min_{\beta} MSE(\beta), \quad MSE(\beta) = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

- **Regularized estimator:**

$$\hat{\beta} = \arg \min_{\beta} J(\beta), \quad J(\beta) = MSE(\beta) + \phi(\beta)$$

- $MSE(\beta)$ = mean-squared prediction error from before
- $\phi(\beta)$ = regularizing function
- **Concept:** Regularizer penalizes β that are “unlikely”
 - Constrains estimate to smaller set of parameters

Two Common Regularizers


- Ridge regression (called L2)

$$\phi(\beta) = \frac{\alpha}{n} \sum_{j=1}^d |\beta_j|^2$$

- LASSO regression (called L1)

$$\phi(\beta) = 2\alpha \sum_{j=1}^d |\beta_j|$$

- Coefficient $\alpha > 0$ determines regularization level
 - Higher $\alpha \Rightarrow$ Higher level of regularization, more constrained
 - Will show how to select α later via cross-validation
 - Scaling factors adjusted to match sklearn convention
- Both penalize large β_j : Tries to make β_j small
 - Will see that L1 also promotes **sparsity**
- Convention: Do not include intercept term β_0
 - In general, no reason to make this term small



L1 and L2 Norm

- Ridge and LASSO Regularization can be written with **norms**

- **Ridge** cost function:

$$J(\boldsymbol{\beta}) = \sum_{i=1}^n (y_i - \hat{y}_i)^2 + \alpha \sum_{j=1}^d |\beta_j|^2 = \|\mathbf{y} - \mathbf{A}\boldsymbol{\beta}\|^2 + \alpha \|\boldsymbol{\beta}\|_2^2$$


- **LASSO** cost function:

$$\begin{aligned} J(\boldsymbol{\beta}) &= \frac{1}{2n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 + \alpha \sum_{j=1}^d |\beta_j| \\ &= \frac{1}{2n} \|\mathbf{y} - \mathbf{A}\boldsymbol{\beta}\|^2 + \alpha \|\boldsymbol{\beta}\|_1 \end{aligned}$$

- $\|\boldsymbol{\beta}\|_1$ = **L1 norm** (pronounced ell-1)

Ridge vs LASSO

- L2 tends to lead to have many “small” coefficients
 - But solutions are not exactly zero
 - Not ideal for feature selection
- L1 tends to lead to more **sparse** solutions
 - Several coefficients are zero



Solving Ridge Regression

- Ridge regression problem: Find β to minimize

$$J(\beta) = \|\mathbf{y} - A\beta\|^2 + \alpha\|\beta\|^2$$

- Solution for given regularization level

$$\beta_{ridge} = (A^T A + \alpha I)^{-1} A^T \mathbf{y}$$

- Set gradient = 0
- Sklearn function for ridge regression:
 - http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.Ridge.html

Solving LASSO Regression

- LASSO cost function:

$$J(\boldsymbol{\beta}) = \frac{1}{2n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 + \alpha \sum_{j=1}^d |\beta_j| = \frac{1}{2n} \|\mathbf{y} - \mathbf{A}\boldsymbol{\beta}\|^2 + \alpha \|\boldsymbol{\beta}\|_1$$

- Because derivative of $|\beta_j|$ is not continuous, there is no closed-form solution.
- Many methods to solve iteratively
 - Least angle regression (LAR), coordinate descent, ADMM
 - However, the cost function is convex \Rightarrow no local minima
 - Beyond the scope of this class
 - See textbook [Hastie2008] for LAR method
 - We will talk about a method later.

Data Scaling

- **Scaling:** Whenever using regularization:
 - Scale each feature and the target to have zero mean and unit variance (or STD)
 - $x_{ij} \rightarrow (x_{i,j} - \bar{x}_j) / \text{STD}(x_{ij})$
 - $y_i \rightarrow (y_i - \bar{y}) / \text{STD}(y_i)$
- After predictor for the scaled data are determined:
 - Derive the equivalent predictor on the original data
- Motivation:
 - Without scaling, the regularization level depends on the data range
 - With mean removal, we do not need the intercept term β_0
 - So that the regularization term is simply a L2 or L1 norm of coefficient vector

Selecting the Regularization Level

- How do we select regularization level α ?
 - Higher $\alpha \Rightarrow$ More constrained / simpler model
 - Lower $\alpha \Rightarrow$ More complex model
- Similar to inverse of model order
- Find α via cross-validation

Pseudo-code

Split in training X_{tr}, y_{tr} and test X_{ts}, y_{ts} .

For α in α_{test} :

- $\hat{\beta} = \text{fit}(X_{tr}, y_{tr}, \alpha)$ // Fit on training data
- $\hat{y}_{ts} = \text{predict}(X_{ts})$ // Predict on test data
- $S[\alpha] = \text{score}(y_{ts}, \hat{y}_{ts})$ // Score on test data

$\hat{\alpha} = \text{argmax } S[\alpha]$ // Select α with highest test score

Summary

Method	Regularizer	Effect on parameters	Solution for Fitting
None	$\phi(\boldsymbol{\beta}) = 0$	Leaves parameters unconstrained	$\hat{\boldsymbol{\beta}} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{y}$
Ridge	$\phi(\boldsymbol{\beta}) = \frac{\alpha}{n} \ \boldsymbol{\beta}\ _2^2$	Makes parameters small Close to zero	$\hat{\boldsymbol{\beta}} = (\mathbf{A}^T \mathbf{A} + \alpha \mathbf{I})^{-1} \mathbf{A}^T \mathbf{y}$
LASSO	$\phi(\boldsymbol{\beta}) = 2\alpha \ \boldsymbol{\beta}\ _1$	Makes parameters sparse. Many coefficients exactly zero	No analytic solution. Need to run an optimizer

- Regularized least squares

$$\hat{\boldsymbol{\beta}} = \arg \min_{\boldsymbol{\beta}} J(\boldsymbol{\beta}), \quad J(\boldsymbol{\beta}) = \frac{1}{n} \|\mathbf{y} - \mathbf{A}\boldsymbol{\beta}\|^2 + \phi(\boldsymbol{\beta})$$

- Whatever you choose for the regularizer:
 - Scale data before training
 - Select regularization level with cross-validation

Outline

- Motivating Example: Feature selection in predicting house prices
- Model selection and regularization
- Housing price prediction with LASSO

LASSO Regression in Python

- Sklearn built in Lasso class
- Easy to use
 - Set alpha
 - Fit on training data
 - Predict and score on test

```
1 from sklearn.linear_model import Lasso
2 from warnings import simplefilter
3 from sklearn.exceptions import ConvergenceWarning
4 simplefilter("ignore", category=ConvergenceWarning)
5
6 # Select alpha
7 alpha = 3e-3
8
9 # Create Lasso object and fit on training data
10 reg = Lasso(alpha=alpha)
11 reg.fit(Xtr1, ytr1)
12
13 # Predict and score on test
14 yhat1 = reg.predict(Xts1)
15 rsq = r2_score(yts1, yhat1)
16
17 print('Test R^2= %f' % rsq)
```

Test R^2= 0.899122

Optimizing Alpha via Cross Validation

- In each fold we:
 - Split data into training and test
 - Fit the scale on the training
 - Transform training and test
 - For each alpha:
 - Fit training and score on test
- Note: Scaling is redone on each fold
 - Ensures scaling is part of the training

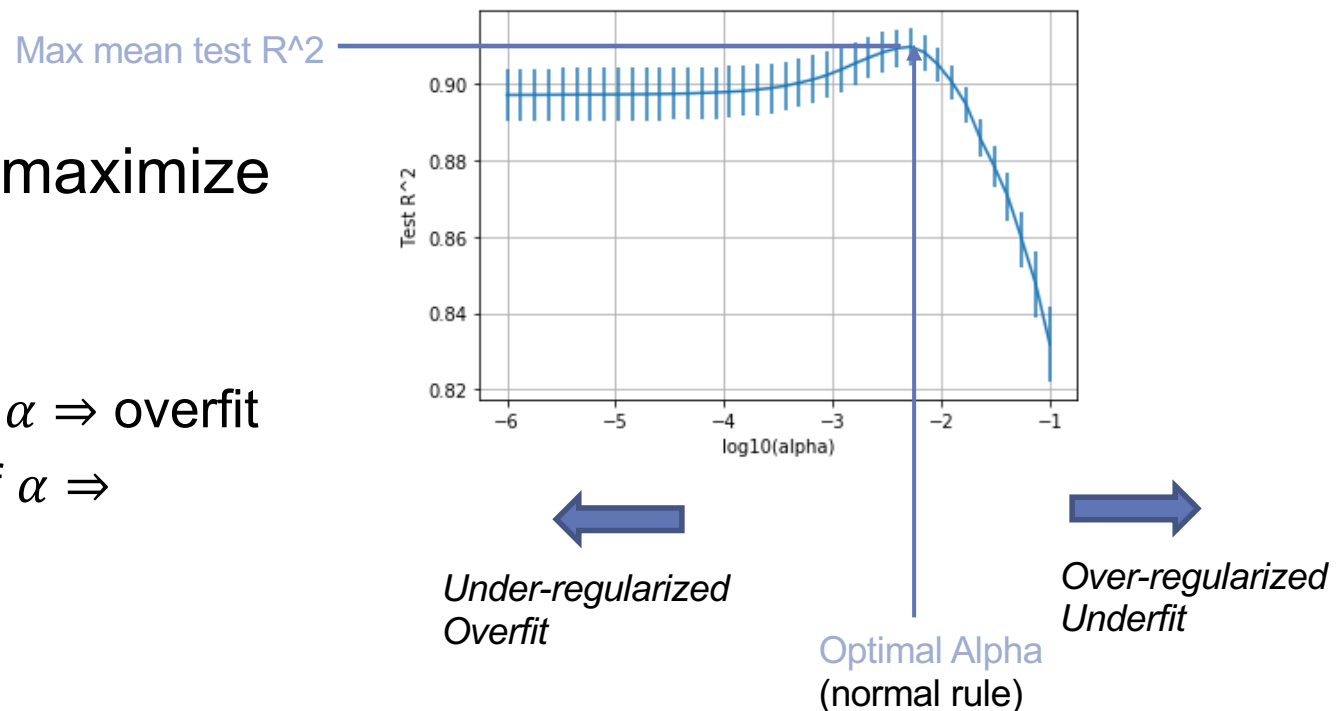
```

10 # Run the cross-validation
11 rsq = np.zeros((nalpha, nfold))
12 for ifold, ind in enumerate(kf.split(X)):
13
14     # Get the training data in the split
15     Itr, Its = ind
16     Xtr = X[Itr,:]
17     ytr = y[Itr]
18     Xts = X[Its,:]
19     yts = y[Its]
20
21     # Fit and transform the data
22     Xtr1 = xscal.fit_transform(Xtr)
23     Xts1 = xscal.transform(Xts)
24     ytr1 = yscal.fit_transform(ytr[:,None])
25     yts1 = yscal.transform(yts[:,None])
26
27     for i, alpha in enumerate(alphas):
28
29         # Fit on the training data
30         reg = Lasso(alpha=alpha)
31         reg.fit(Xtr1, ytr1)
32
33         # Score on the test data
34         yhat1 = reg.predict(Xts1)
35         rsq[i, ifold] = r2_score(yts1, yhat1)
36
37     print('Fold = %d' % ifold)
38
39 # Compute mean and SE
40 rsq_lasso_mean = np.mean(rsq, axis=1)
41 rsq_lasso_se = np.std(rsq, axis=1) / np.sqrt(nfold-1)

```

Cross Validation: Normal Rule

- Select alpha to maximize mean test R^2
 - Normal rule
 - Lower values of $\alpha \Rightarrow$ overfit
 - Higher values of $\alpha \Rightarrow$ underfit

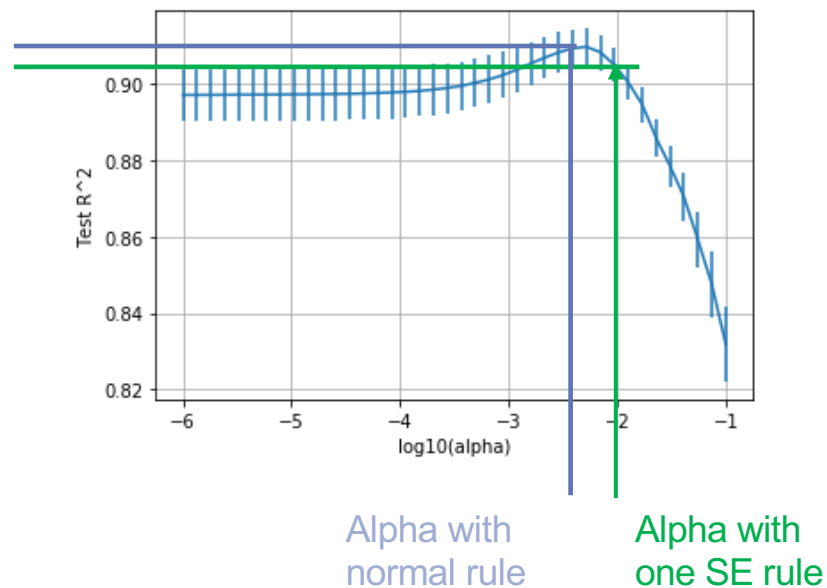


Alpha optimal (normal rule) = $5.2233e-03$
 Mean test R^2 (normal rule) = 0.910

Cross Validation: One SE Rule

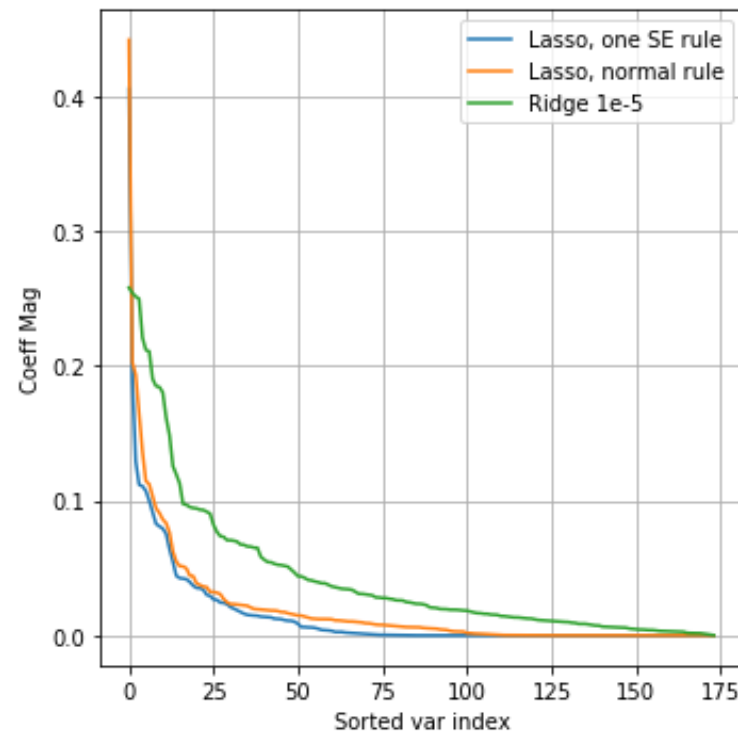

Max mean test R^2
Max mean test R^2 -one SE

- Can also use one SE rule:
 - Selects a higher regularized model
 - More sparse solution



Sparsity in the Coefficients

- Adding L1 regularization:
 - Makes coefficient smaller
 - Many coefficients approx. 0




Most Important Variables

- Right table: Variables with 10 large coefficient magnitudes
- Minimally regularized (Ridge) has:
 - Variables that are highly correlated
 - Ex: GrLivArea and 2ndFlrSF
 - Several large variables
- Lasso:
 - Reduces correlated variables
 - Selects GrLivArea alone
 - Gives the variables more importance
 - GrLivArea coefficient is much higher

Ridge		Lasso	
GrLivArea	0.29	GrLivArea	0.42
2ndFlrSF	0.26	OverallQual	0.18
KitchenQual_Gd	-0.21	KitchenQual_TA	-0.17
KitchenQual_TA	-0.20	KitchenQual_Gd	-0.16
LotArea	0.18	YearBuilt	0.13
YearBuilt	0.16	BsmtFinSF1	0.12
OverallQual	0.16	Neighborhood_NoRidge	0.09
ExterQual_Gd	-0.15	OverallCond	0.09
Exterior2nd_VinylSd	0.15	TotalBsmtSF	0.09
ExterQual_TA	-0.15	LotArea	0.08

Ridge Vs. Lasso

- Can optimize alpha for both regularizers
- Optimal mean test R^2 is better for LASSO
- Offers better feature selection



Optimal R^2 Lasso: 0.907283
Optimal R^2 Ridge: 0.897498

Lasso Path

- Plot of coefficients vs. alpha
- For large alpha:
 - All coefficients are zero
- As alpha is decreased:
 - One coefficient is activated at a time
 - Indicates an ordering of importance

