

# LECTURE 8: SUPPORT VECTOR MACHINES

---

Ehsan Aryafar

earyafar@pdx.edu

<http://web.cecs.pdx.edu/~aryafare/ML.html>

# Recall: Gradient Defined

- Consider scalar-valued function  $f(\mathbf{w})$
- Vector input  $\mathbf{w}$ . Then gradient is:


$$\nabla_{\mathbf{w}} f(\mathbf{w}) = \begin{bmatrix} \partial f(\mathbf{w}) / \partial w_1 \\ \vdots \\ \partial f(\mathbf{w}) / \partial w_N \end{bmatrix}$$

- Matrix input  $\mathbf{W}$ , size  $M \times N$ . Then gradient is:


$$\nabla_{\mathbf{W}} f(\mathbf{W}) = \begin{bmatrix} \partial f(\mathbf{W}) / \partial W_{11} & \cdots & \partial f(\mathbf{W}) / \partial W_{1N} \\ \vdots & \vdots & \vdots \\ \partial f(\mathbf{W}) / \partial W_{M1} & \cdots & \partial f(\mathbf{W}) / \partial W_{MN} \end{bmatrix}$$

- Gradient is same size as the argument!

# Recall: Gradient Descent Illustrated



- $M = 1$



- $M = 2$

# Recall: Convex Sets


- **Definition:** A set  $X$  is **convex** if for any  $x, y \in X$ ,

$$tx + (1 - t)y \in X \text{ for all } t \in [0,1]$$


- Any line between two points remains in the set.
- Examples:
  - Square, circle, ellipse
  - $\{x \mid Ax \leq b\}$  for any matrix  $A$  and vector  $b$

# Recall: Convex Set Visualized

- Convex




- Not convex




# Recall: Convex Functions

- A real-valued function  $f(x)$  is **convex** if:
  - Its domain is a convex set, and
  - For all  $x, y$  and  $t \in [0,1]$ :
$$f(tx + (1 - t)y) \leq tf(x) + (1 - t)f(y)$$



# Recall: Global Minima and Convex Function

- **Theorem:** If  $f(w)$  is convex and  $w$  is a local minima, then  $w$  is a global minima
- Implication for optimization:
  - Gradient descent only converges to local minima
  - In general, cannot guarantee optimality
  - Depends on initial condition
  - But, for convex functions can always obtain optimal



# Learning Objectives

- Interpret weights in linear classification of images
- Describe why linear classification for images does not work
- Define the margin in linear classification
- Describe the SVM classification problem.
- Describe a kernel SVM problem for non-linear classification
- Implement SVM classifiers in python
- Select SVM parameters from cross-validation



# Outline

- Motivating example: Recognizing handwritten digits
  - Why logistic regression doesn't work well.
- Maximum margin classifiers
- Support vector machines
- Kernel trick

# MNIST Digit Classification

HANDWRITING SAMPLE FORM

NAME [REDACTED] DATE 8-3-89 CITY MINDEN CITY STATE MI ZIP 48456

This sample of handwriting is being collected for use in testing computer recognition of hand printed numbers and letters. Please print the following characters in the boxes that appear below.

0 1 2 3 4 5 6 7 8 9      0 1 2 3 4 5 6 7 8 9      0 1 2 3 4 5 6 7 8 9

0123456789    0123456789    0123456789

87      701      3752      \*0759      960941

87    701    3752    \*0759    960941

158      4586      32123      832656      82

158    4586    32123    832656    82

7481      80539      419219      67      904

7481    80539    419219    67    904

61738      729658      75      390      5716

61738    729658    75    390    5716

From Patrick J. Grother, NIST Special Database, 1995

- Problem: Recognize hand-written digits
- Original problem:
  - Census forms
  - Automated processing
- Classic machine learning problem
- Benchmark

# A Widely-Used Benchmark

## Classifiers [\[ edit \]](#)

This is a table of some of the [machine learning](#) methods used on the database and their error rates, by type of classifier:

Type ↕	Classifier ↕	Distortion ↕	Preprocessing ↕	Error rate (%) ↕
Linear classifier	<a href="#">Pairwise linear classifier</a>	None	Deskewing	7.6 <sup>[9]</sup>
<a href="#">K-Nearest Neighbors</a>	K-NN with non-linear deformation (P2DHMDM)	None	Shiftable edges	0.52 <sup>[14]</sup>
<a href="#">Boosted Stumps</a>	Product of stumps on <a href="#">Haar features</a>	None	Haar features	0.87 <sup>[15]</sup>
Non-Linear Classifier	40 PCA + quadratic classifier	None	None	3.3 <sup>[9]</sup>
<a href="#">Support vector machine</a>	Virtual SVM, deg-9 poly, 2-pixel jittered	None	Deskewing	0.56 <sup>[16]</sup>
<a href="#">Neural network</a>	2-layer 784-800-10	None	None	1.6 <sup>[17]</sup>
<a href="#">Neural network</a>	2-layer 784-800-10	<a href="#">elastic distortions</a>	None	0.7 <sup>[17]</sup>
<a href="#">Deep neural network</a>	6-layer 784-2500-2000-1500-1000-500-10	<a href="#">elastic distortions</a>	None	0.35 <sup>[18]</sup>
<a href="#">Convolutional neural network</a>	Committee of 35 conv. net, 1-20-P-40-P-150-10	<a href="#">elastic distortions</a>	Width normalizations	0.23 <sup>[8]</sup>

- We will look at SVM today
- Not the best algorithm
- But quite good
- ...and illustrates the main points



# Downloading MNIST

```
import tensorflow as tf

(Xtr,ytr),(Xts,yts) = tf.keras.datasets.mnist.load_data()

print('Xtr shape: %s' % str(Xtr.shape))
print('Xts shape: %s' % str(Xts.shape))

ntr = Xtr.shape[0]
nts = Xts.shape[0]
nrow = Xtr.shape[1]
ncol = Xtr.shape[2]
```

```
Xtr shape: (60000, 28, 28)
Xts shape: (10000, 28, 28)
```

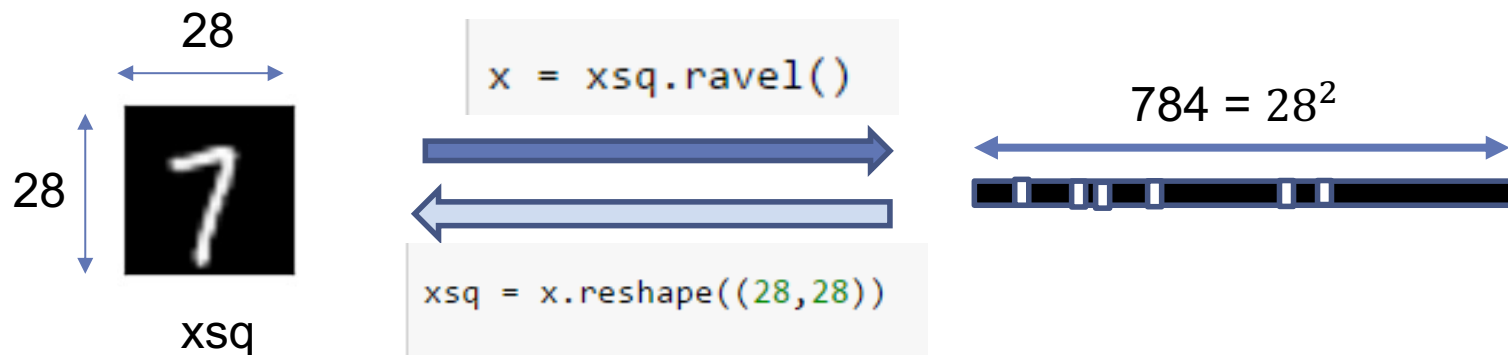
- MNIST data is available in many sources
  - Note: It has been removed from sklearn
- Tensorflow version:
  - 60000 training samples
  - 10000 test samples
- Each sample is a 28 x 28 image
- Grayscale: Pixel values  $\in \{0, 1, \dots, 255\}$ 
  - 0 = Black and
  - 255 = White

# Matrix and Vector Representation

- For this demo, we reshape data from  $N \times 28 \times 28$  to  $N \times 784$
- But, you can easily go back and forth
- Also, scale the pixel values from -1 to 1

```
npix = nrow*ncol
Xtr = 2*(Xtr/255 - 0.5)
Xtr = Xtr.reshape((ntr,npix))

Xts = 2*(Xts/255 - 0.5)
Xts = Xts.reshape((nts,npix))
```



$$S = \text{Mat}(x) = \begin{bmatrix} s_{11} & \cdots & s_{1,28} \\ \vdots & \vdots & \vdots \\ s_{28,1} & \cdots & s_{28,28} \end{bmatrix}$$

$$x = \text{vec}(S) = [x_1 \quad \cdots \quad x_{784}]$$

# Displaying Images in Python



4 random images in the dataset

We want to classify each digit

A human can classify these easily

Getting a computer to do is harder

```
def plt_digit(x):  
    nrow = 28  
    ncol = 28  
    xsq = x.reshape((nrow,ncol))  
    plt.imshow(xsq, cmap='Greys_r')  
    plt.xticks([])  
    plt.yticks([])  
  
    # Convert data to a matrix  
    X = mnist.data  
    y = mnist.target  
  
    # Select random digits  
    nplt = 4  
    nsamp = X.shape[0]  
    Iperm = np.random.permutation(nsamp)  
  
    # Plot the images using the subplot command  
    for i in range(nplt):  
        ind = Iperm[i]  
        plt.subplot(1,nplt,i+1)  
        plt_digit(X[ind,:])
```

Key command

Sample permutation is necessary for this dataset, as the original data is ordered by digits

# Try a Logistic Classifier

```
ntr1 = 5000
Xtr1 = Xtr[Iperm[:ntr1],:]
ytr1 = ytr[Iperm[:ntr1]]
```

- Train on 5000 samples
  - To reduce training time.
  - In practice want to train with ~40k
- Select correct solver (sag)
  - Others can be very slow. Even this will take minutes

```
from sklearn import linear_model
logreg = linear_model.LogisticRegression(verbose=10, solver='sag',\
                                         max_iter=500)
logreg.fit(Xtr1,ytr1)
```

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
convergence after 408 epochs took 64 seconds
```

```
[Parallel(n_jobs=1)]: Done    1 out of    1 | elapsed:   1.1min remaining:    0.0s
[Parallel(n_jobs=1)]: Done    1 out of    1 | elapsed:   1.1min finished
```

# Performance

- Accuracy = 89%. Very bad
- Some of the errors seem like they should have been easy to spot
- What went wrong?

```
nts1 = 5000
Iperm_ts = np.random.permutation(nts)
Xts1 = Xts[Iperm_ts[:nts1],:]
yts1 = yts[Iperm_ts[:nts1]]
yhat = logreg.predict(Xts1)
acc = np.mean(yhat == yts1)
print('Accuracy = {0:f}'.format(acc))
```

Accuracy = 0.891000

true=7 est=0



true=8 est=2



true=3 est=5




true=9 est=5





# Recap: Logistic Classifier




- Each **logit**  $z_k = \mathbf{w}_k^T \mathbf{x}$  = inner product with **weight**  $\mathbf{w}_k$  with digit  $\mathbf{x}$ ,  $k = 0, \dots, 9$
- Will select  $\hat{y} = \arg \max_k P(y = k|x) = \arg \max_k z_k$ 
  - Output  $z_k$  which is largest
- When is  $z_k$  large?

# Interpreting the Logistic Classifier Weights

- A logit  $z_k = \mathbf{w}_k^T \mathbf{x}$  is high when there is high **overlap** between  $\mathbf{w}_k$  with digit  $x$ 
  - Visualize each weight as an image
  - Suppose pixels are 0 or 1
  - $z_k = \mathbf{w}_k^T \mathbf{x} = \sum_i w_{ki} x_i =$  number of pixels that overlap with  $\mathbf{w}_k$  and  $x$
- **Conclusion:** Small variations in digits can cause low overlap


Weight for digit “0”

$\mathbf{w}_0$



Digit

$x$



Overlap **high**  $\Rightarrow x$  **is** digit 0

Overlap **low**  $\Rightarrow x$  **not** digit 0


# Example with Actual Digits

- Take weight  $w$  from a random digit “2”
- Inner products  $z = w^T x$  are only slightly higher for other digits “2”
- Cannot tell which digit is correct from the inner product  $z = w^T x$



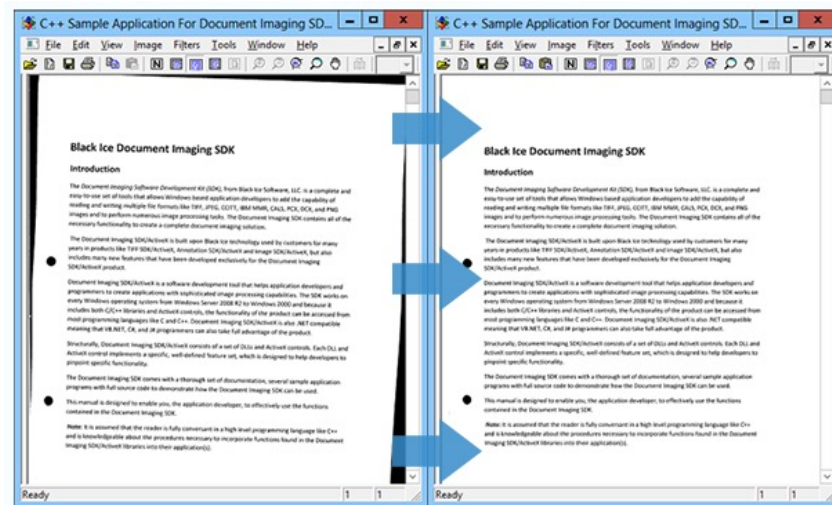
# Visualizing the Weights

- Optimized weights of the classifier
- Blurry versions of image to try to capture rotations, translations, ...



# Problems with Logistic Classifier

- Linear weighting cannot capture many deformities in image
  - Rotations
  - Translations (movement)
  - Variations in relative size of digit components
- Can be improved with preprocessing
  - E.g. deskewing, contrast normalization, many methods
- Is there a better classifier?




# Outline

- Motivating example: Recognizing handwritten digits
  - Why logistic regression doesn't work well.
- Maximum margin classifiers
- Support vector machines
- Kernel trick


# Non-Uniqueness of Separating Plane

- Linearly separable data:
  - Can find a separating hyper-plane as a linear classifier.
- Separating hyper-plane is not unique
  - Fig. on right: Many separating planes
- Which one is optimal?



# Hyperplane Basics


- Linear function:  $f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b, \mathbf{x} \in R^d$
- Hyperplane in d-dimensional:  $f(\mathbf{x}) = 0$
- Parameters:
  - Weight  $\mathbf{w}$  and bias  $b$
  - Unique up to scaling:
  - $(b, \mathbf{w})$  and  $(\alpha b, \alpha \mathbf{w})$  define the same plane.
  - For unique definition, we can require  $\|\mathbf{w}\|=1$ .
- Distance of any point  $\mathbf{x}$  to the hyperplane:
  - $d = f(\mathbf{x})/\|\mathbf{w}\|$ , where  $f(\mathbf{x}) = b + \mathbf{w}^T \mathbf{x}$ .
  - See ESL Sec. 4.5.
  - ESL: Hastie, Tibshirani, Friedman, “The Elements of Statistical Learning”. 2<sup>nd</sup> Ed. Springer.







# Linear Separability and Margin

- Given training data  $(\mathbf{x}_i, y_i), i = 1, \dots, N$ 
  - Binary class label:  $y_i = \pm 1$
- Suppose it is separable with parameters  $(\mathbf{w}, b)$
- There must exist a  $\gamma > 0$  s.t.:
  - $b + w_1x_{i1} + \dots w_dx_{id} > \gamma$  when  $y_i = 1$
  - $b + w_1x_{i1} + \dots w_dx_{id} < -\gamma$  when  $y_i = -1$
- Single equation form:
 
$$y_i(b + w_1x_{i1} + \dots w_dx_{id}) > \gamma \text{ for all } i = 1, \dots, N$$
- **Margin:**  $m = \frac{\gamma}{\|\mathbf{w}\|}$  : minimal distance of a sample to the plane
  - $\gamma$  is the minimum value satisfying the above constraints



# Which separating plane is better ?



From Fig. 9.2 and Fig. 9.3 in ISL.

# Maximum Margin Classifier

- For the classifier to be more robust to noise, we want to maximize the margin!

- Define maximum margin classifier optimization problem


$$\begin{array}{ll} \max_{w, \gamma} \gamma & \longleftarrow \text{Maximizes the margin} \\ \text{Such that } y_i(b + \mathbf{w}^T \mathbf{x}) \geq \gamma \text{ for all } i & \longleftarrow \text{Ensures all points are correctly classified} \\ \sum_{j=1}^d w_j^2 \leq 1 & \longleftarrow \text{Scaling on weights} \end{array}$$

- Called a constrained optimization problem

- Objective function and constraints
- More on this later.

- See closed form solution in Sec. 4.5.2 in ESL. Note notation difference.

# Visualizing Maximum Margin Classifier



- Fig. 9.3 of ISL
- Margin determined by closest points to the line
  - The maximal margin hyperplane represents the mid-line of the **widest “slab”** that we can insert between two classes
- In this figure, there are 3 points at the margin

ISL: James, Witten, Hastie, Tibshirani, An Introduction to Statistical Learning, Springer. 2013.

# Problems with MM classifier

- Data is often not perfectly separable
  - You cannot talk about margin
  - Only want to correctly separate most points
- MM classifier is not robust
  - A single sample can radically change line
  - Suggests generalization errors may not be good




Fig. 9.4





Fig. 9.5

# Outline

- Motivating example: Recognizing handwritten digits
  - Why logistic regression doesn't work well.
- Maximum margin classifiers
- Support vector machines
- Kernel trick


# Support Vector Machine

- Support Vector Machine (SVM)
  - Vladimir Vapnik, 1963
  - But became widely-used with kernel trick, 1993
  - More on this later
- Got best results on character recognition
- Key idea: Allow “slack” in the classification
  - Support vector classifier (SVC): Directly use raw features. Good when the original feature space is roughly linearly separable
  - Support vector machine (SVM): Map the raw features to some other domain through a kernel function



# Hinge Loss

- Fix  $\gamma = 1$
- Want ideally:  $y_i(\mathbf{w}^T \mathbf{x} + b) \geq 1$  for all samples  $i$ 
  - Equivalently,  $y_i z_i \geq 1$ ,  $z_i = b + \mathbf{w}^T \mathbf{x}$
  - Note that  $y_i$  is + or - one
- But perfect separation may not be possible
- Define **hinge loss** or **soft margin**:
  - $L_i(\mathbf{w}, b) = \max(0, 1 - y_i z_i)$
- Starts to increase as sample is misclassified:
  - $y_i z_i \geq 1 \Rightarrow$  Sample meets margin target,  $L_i(\mathbf{w}) = 0$
  - $y_i z_i \in [0, 1) \Rightarrow$  Sample margin too small, small loss
  - $y_i z_i \leq 0 \Rightarrow$  Sample misclassified, large loss





# SVM Optimization

- Given data  $(\mathbf{x}_i, y_i)$
- Optimization  $\min_{\mathbf{w}, b} J(\mathbf{w}, b)$

$$J(\mathbf{w}, b) = C \sum_{i=1}^N \max(0, 1 - y_i(\mathbf{w}^T \mathbf{x}_i + b)) + \frac{1}{2} \|\mathbf{w}\|^2$$

C controls final  
margin

Hinge loss term  
Attempts to reduce  
Misclassifications

margin =  $1/\|\mathbf{w}\|$

- Constant  $C > 0$  will be discussed below
- Note: ISL book uses different naming conventions.
  - We have followed convention in sklearn

# Alternate Form of SVM Optimization (Constrained Optimization Format)

- Equivalent optimization:

$$\min J_1(\mathbf{w}, b, \epsilon), \quad J_1(\mathbf{w}, b, \epsilon) = C \sum_{i=1}^N \epsilon_i + \frac{1}{2} \|\mathbf{w}\|^2$$

- Subject to constraints:

$$y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1 - \epsilon_i \text{ for all } i = 1, \dots, N$$


- $\epsilon_i$  = amount sample  $i$  misses margin target
- Sometimes written as  $J_1(\mathbf{w}, b, \epsilon) = C \|\epsilon\|_1 + \frac{1}{2} \|\mathbf{w}\|^2$ 
  - $\|\epsilon\|_1 = \sum_{i=1}^N \epsilon_i$  called the “one-norm”
  - Generally one-norm would have absolute sign over  $\epsilon_i$ .
  - But in this case, when the constraint is met,  $\epsilon_i \geq 0$ .

# Interpreting Parameters

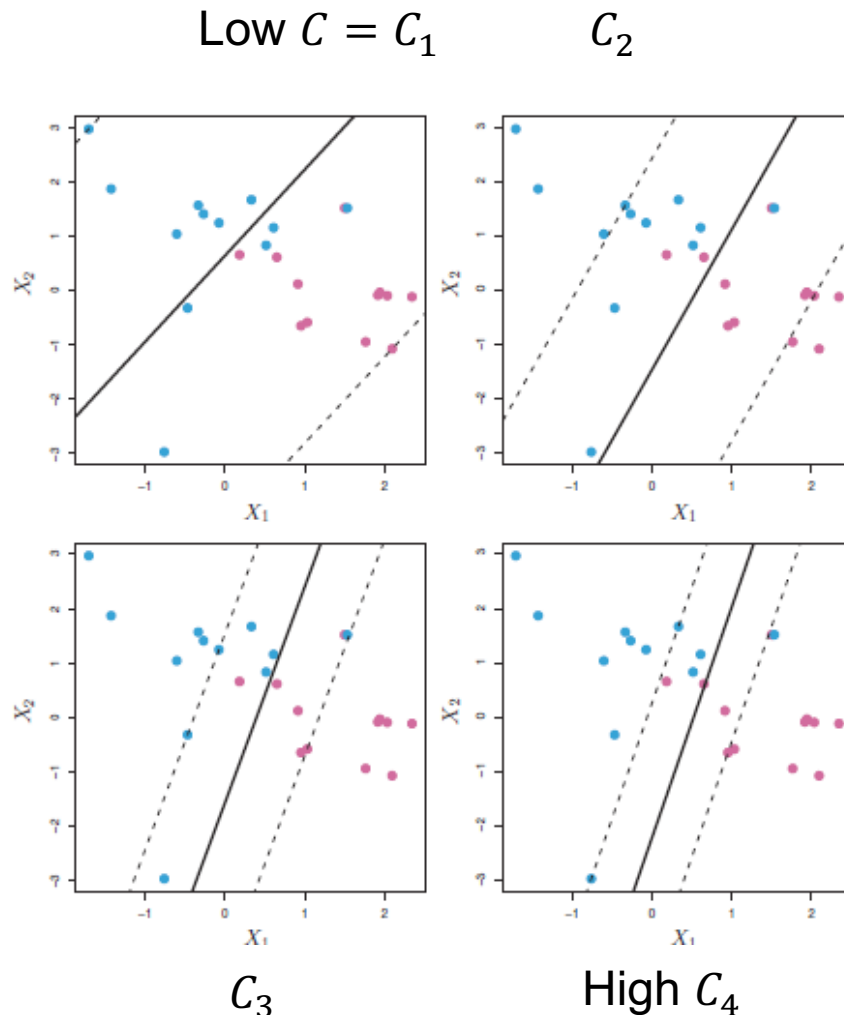
- Margin is  $1/\|\mathbf{w}\|$
- Parameter  $\epsilon_i$  called the slack variable
  - $\epsilon_i = 0 \Rightarrow$  Sample on correct side of margin
  - $0 \leq \epsilon_i < 1 \Rightarrow$  Sample violates the margin (are inside the margin)
  - $\epsilon_i \geq 1 \Rightarrow$  Sample misclassified (wrong side of hyperplane)
- Parameter  $C$  (Discussed Soon):
  - Balance between first term (violations) and second term (inverse of margin)
  - $C$  large: Forces minimum number of violations, but small margin.
    - Highly fit to data. Low bias, higher variance
  - $C$  small: Enables more samples violations, but large margin.
    - Higher bias, lower variance
  - Found by cross-validation

# Support Vectors

- **Support vectors:** Samples that either:
  - Are exactly on margin:  $y_i(\mathbf{w}^T \mathbf{x}_i + b) = 1$
  - Or, on wrong side of margin:  $y_i(\mathbf{w}^T \mathbf{x}_i + b) \leq 1$
- Changing samples that are not SVs
  - Does not change solution
  - Provides robustness



# Illustrating Effect of $C$




- Fig. 9.7 of ISL
  - Note:  $C$  has opposite meaning in ISL than python
  - Here, we use python meaning
- Low  $C$ :
  - Leads to large margin
  - But allow many violations of margin.
  - Many more SVs
  - Reduces variance or increases bias by using more samples
- Large  $C$ :
  - Leads to small margin
  - Reduce number of violations, and fewer SVs.
  - Highly fit to data. Low bias, higher variance
  - More chance to overfit

# Relation to Logistic Regression

- Logistic regression also minimizes a loss function:

$$J(\mathbf{w}, b) = \sum_{i=1}^N L_i(\mathbf{w}, b),$$

$$L_i(\mathbf{w}, b) = \ln P(y_i | \mathbf{x}_i) = -\ln(1 + e^{-y_i z_i})$$




# Outline

- Motivating example: Recognizing handwritten digits
  - Why logistic regression doesn't work well.
- Maximum margin classifiers
- Support vector machines
- **Kernel trick**

# The Kernel Function

- Kernel function:
  - Function  $K(x_i, x)$
  - Key function for SVMs and kernel classifiers
  - Measures “similarity” between new sample  $x$  and training sample  $x_i$
- Typical property
  - $x_i, x$  close  $\Rightarrow K(x_i, x)$  maximum value
  - $x_i, x$  far  $\Rightarrow K(x_i, x) \approx 0$






# Common Kernels

- Linear SVM:

- $K(x_i, x) = x_i^T x = \|x_i\| \|x\| \cos \theta$
- Maximum when angle between vectors is small




- Radial basis function:

$$K(x_i, x) = \exp[-\gamma \|x - x_i\|^2]$$

- $1/\gamma$  indicates width of kernel

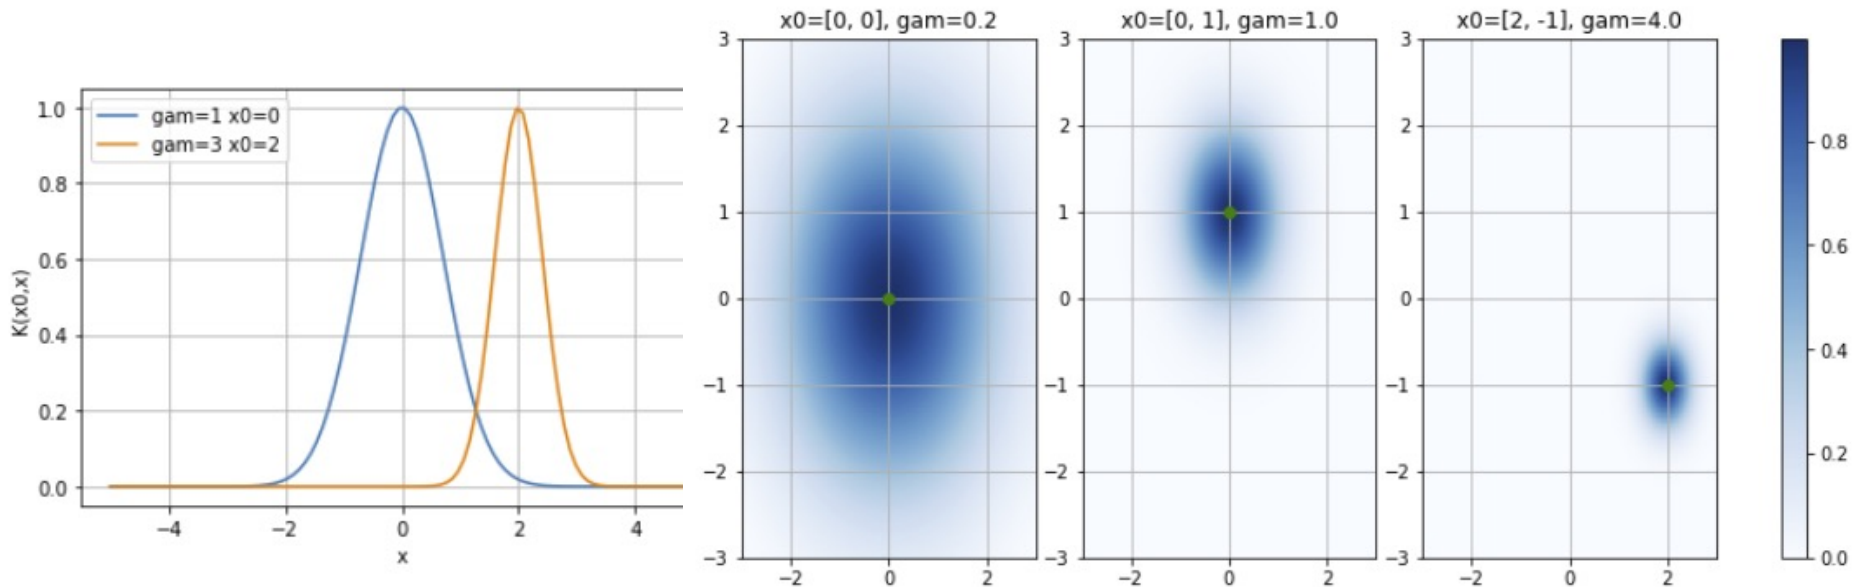
- Polynomial kernel:  $K(x_i, x) = |x_i^T x|^d$

- Inner product to the power of  $d$ !
- Typically  $d=2$



# RBF Kernel Examples

- RBF kernel:  $K(x_0, x) = \exp[-\gamma \|x - x_0\|^2]$ 
  - Peak value of 1 at  $x = x_0$
  - Decay with a rate of  $\frac{1}{\gamma}$
  - Width  $\propto \frac{1}{\gamma}$



RBFs in 1D

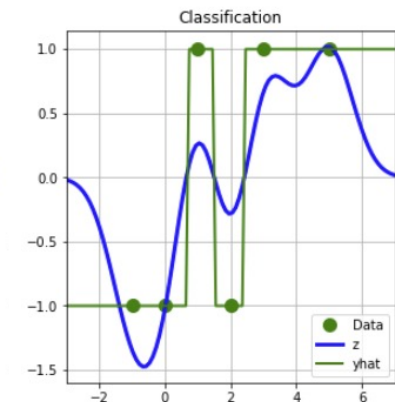
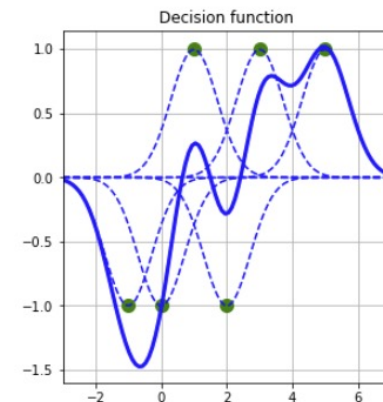
# Kernel Classifier

- Given:
  - Training data  $(x_i, y_i)$  with binary labels  $y_i = \pm 1$
  - Kernel  $K(x_i, x)$
- To classify a new point  $x$ :
  - Decision function:  $z = \sum_{i=1}^n y_i K(x_i, x)$
  - Classify:  $\hat{y} = \text{sign}(z)$
- Idea:
  - $z$  is large positive when  $x$  is close to samples  $x_i$  with  $y_i = 1$
  - $z$  is large negative when  $x$  is close to samples  $x_i$  with  $y_i = -1$
- Kernel classifiers are a subject on their own
  - We just mention them here to explain connection to SVMs

# Example in 1D

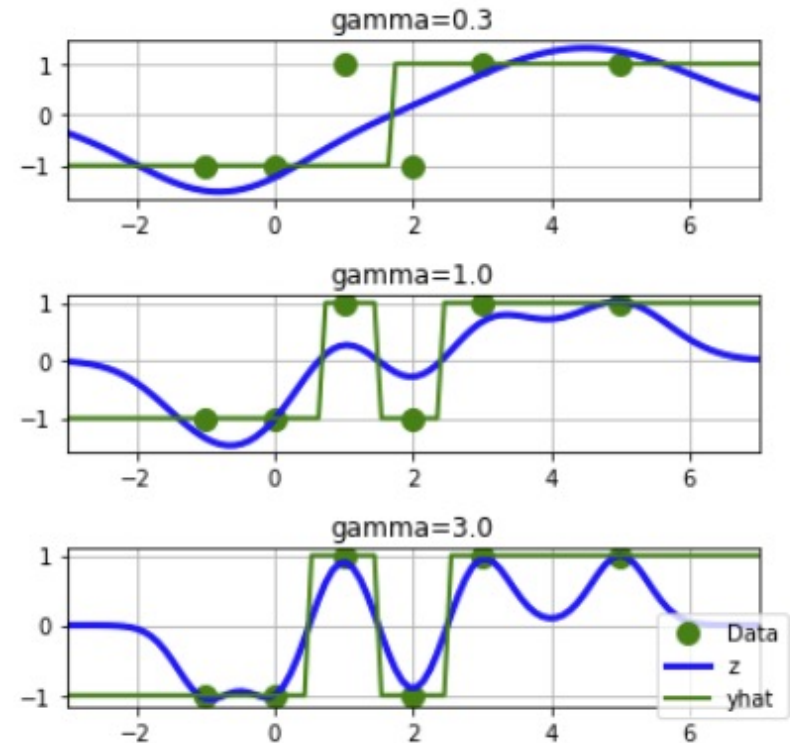
- Example data with 6 points  $(x_i, y_i)$ 
  - RBF kernel:  $K(x_i, x) = e^{-\gamma(x_i - x)^2}$ ,  $\gamma = 1$
- Decision function:
  - $z = \sum_{i=1}^n y_i K(x_i, x)$
  - Sum of bell curves
  - Positive when near positive samples
  - Negative when near negative samples
- Classification:
  - $\hat{y} = \text{sign}(z)$

$i$	1	2	3	4	5	6
$x_i$	-1	0	1	2	3	5
$y_i$	-1	-1	1	-1	1	1



# Effect of Gamma


- Same data as before
- RBF kernel:  $K(x_i, x) = e^{-\gamma(x_i - x)^2}$
- As  $\gamma$  increases:
  - Decision function  $z \approx y_i$  when  $x = x_i$
  - Classifier fits training data better
  - Classification region more complex
- As a classifier, higher  $\gamma$  results in:
  - Lower bias error (fits training data)
  - But, higher variance error
  - Overfitting



# SVMs with Non-Linear Transformations

- Non-linear transformation:
  - Replace  $x$  with  $\phi(x)$
  - Enables more rich, non-linear classification
  - Examples: polynomial classification

$$\phi(x) = [1, x, x^2, \dots, x^{d-1}]$$



- Tries to find separation in a **feature** space (e.g., classification in the picture)
  - You can do this with any classifier (we have already done this)
- **Kernel trick** in SVMs:
  - Makes applying non-linear transformations easy

# SVM with the Transformation

- Consider SVM model with  $x$  replaced by  $\phi(x)$
- Minimize SVM cost function as before (i.e. Hinge loss + inverse margin)
- **Theorem:** The optimal weight is of the form (linear):

$$\mathbf{w} = \sum_{i=1}^N \alpha_i y_i \phi(\mathbf{x}_i)$$

- $\alpha_i \geq 0$  for all  $i$
- $\alpha_i > 0$  if and only if sample  $i$  is a support vector
- Will show this fact later using results in constrained optimization
- **Consequence:** The linear discriminant on any other sample  $x$  is:

$$z = b + \mathbf{w}^T \phi(\mathbf{x}) = b + \sum_{i=1}^N \alpha_i y_i \phi(\mathbf{x}_i)^T \phi(\mathbf{x})$$

$K(\mathbf{x}_i, \mathbf{x}) = \text{"kernel"}$

# Kernel Form of the SVM Classifier

- SVM classifier can be written with the kernel  $K(\mathbf{x}_i, \mathbf{x})$  and values  $\alpha_i \geq 0$ :

$$z = b + \sum_{i=1}^N \alpha_i y_i K(\mathbf{x}_i, \mathbf{x}),$$

← Decision function

$$\hat{y} = \text{sign}(z) = \begin{cases} 1 & \text{if } z > 0 \\ -1 & \text{if } z < 0 \end{cases}$$

← Classification decision

- **Key point:** SVM classifier is approximately Kernel classifier
- But there are two differences:
  - introduction of weights  $\alpha_i \geq 0$  on the samples (the weights are only non-zero on the SVs)
  - A bias term  $b$  (can be positive or negative)



# “Kernel Trick” and Dual Parameterization

- Kernel form of SVM classifier (previous slide):


$$z = b + \sum_{i=1}^N \alpha_i y_i K(\mathbf{x}_i, \mathbf{x}),$$
$$\hat{y} = \text{sign}(z)$$

- Dual parameters:  $\alpha_i \geq 0, i = 1, \dots, N$ 
  - Problem based on  $\alpha_i$  parameters
  - Called the dual parameters due to constrained optimization – see next section
- Kernel trick:
  - Directly solve the parameters  $\alpha$  instead of the weights  $\mathbf{w}$
  - Can show that the optimization only needs the kernel  $K(\mathbf{x}_i, \mathbf{x})$
  - Does not need to explicitly use  $\phi(\mathbf{x})$

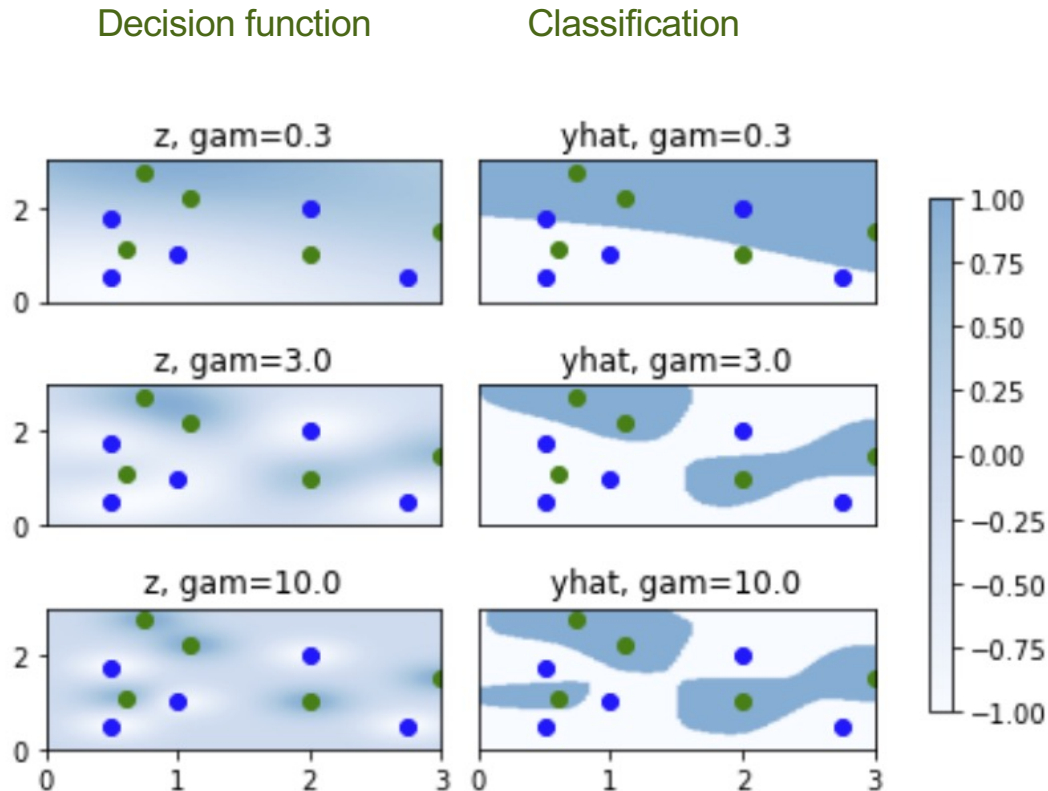
# SVM Example in 1D

$i$	1	2	3	4	5	6
$x_i$	-1	0	1	2	3	5
$y_i$	-1	-1	1	-1	1	1

- Same data as in the Kernel classifier example
- Fit SVM with RBF with different  $\gamma$
- Similar trends as kernel classifier: As  $\gamma$  increases
  - $z$  “fits” data  $(x_i, y_i)$  closer
  - Leads to more complex decision regions.
  - Enables nonlinear decision regions




# Example in 2D



- Example:
  - 10 data points with binary labels
  - Fit SVM with  $C = 1$  and RBF
  - $\gamma = 0.3, 3$  and  $10$
- Plot:
  - $z$  = linear discriminant
  - $\hat{y} = \text{sign}(z)$  = classification decision
- Observe: As  $\gamma$  increases
  - Fits training data better
  - More complex decision region

# Parameter Selection

- For SVMs with RBFs we need to select:
  - Parameter  $C > 0$  in the loss function
  - Kernel width  $\gamma > 0$
- Higher  $C$  or  $\gamma$ 
  - Fewer SVs
  - Classifiers averages over smaller set
  - Lower bias, but higher variance
- Typically select via cross-validation
  - Try out different  $(C, \gamma)$  pairs
  - Find which one provides highest accuracy on test set
- Python can automatically do grid search



<http://peekaboo-vision.blogspot.com/2010/09/mnist-for-ever.html>

# Multi-Class SVMs

- Suppose there are  $K$  classes
- One-vs-one:
  - Train  $\binom{K}{2}$  SVMs for each pair of classes
  - Test sample assigned to class that wins “majority of votes”
  - Best results but very slow
- One-vs-rest:
  - Train  $K$  SVMs: train each class  $k$  against all other classes
  - Pick class with highest  $z_k$
- Sklearn has both options

# MNIST Results

- Run classifier
- Very slow
  - Several minutes for 40,000 samples
  - Slow in training and test
  - Major drawback of SVM
- Accuracy  $\approx 0.984$ 
  - Much better than logistic regression
- Can get better with:
  - pre-processing
  - More training data
  - Optimal parameter selection

```
from sklearn import svm

# Create a classifier: a support vector classifier
svc = svm.SVC(probability=False, kernel="rbf", C=2.8, gamma=.0073, verbose=10)
```

```
svc.fit(Xtr,ytr)
```

```
[LibSVM]
```

```
SVC(C=2.8, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape=None, degree=3, gamma=0.0073, kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=10)
```

```
yhat1 = svc.predict(Xts)
acc = np.mean(yhat1 == yts)
print('Accuracy = {0:f}'.format(acc))
```

```
Accuracy = 0.984000
```

# MNIST Errors

- Some of the error are hard even for a human

true=7 est=9



true=8 est=5



true=5 est=6



true=9 est=4



# What you should know

- Interpret weights in linear classification of images (logistic regression): Match filters
- Understand the margin in linear classification and maximum margin classifier
- SVM classifier: Allow violation of margin by introducing slack variables (More robust than linear classifier)
- Extend to nonlinear classifier by feature transformation: SVM with nonlinear kernels
- Select SVM parameters from cross-validation