

LECTURE 10: PRINCIPAL COMPONENT ANALYSIS (PCA)

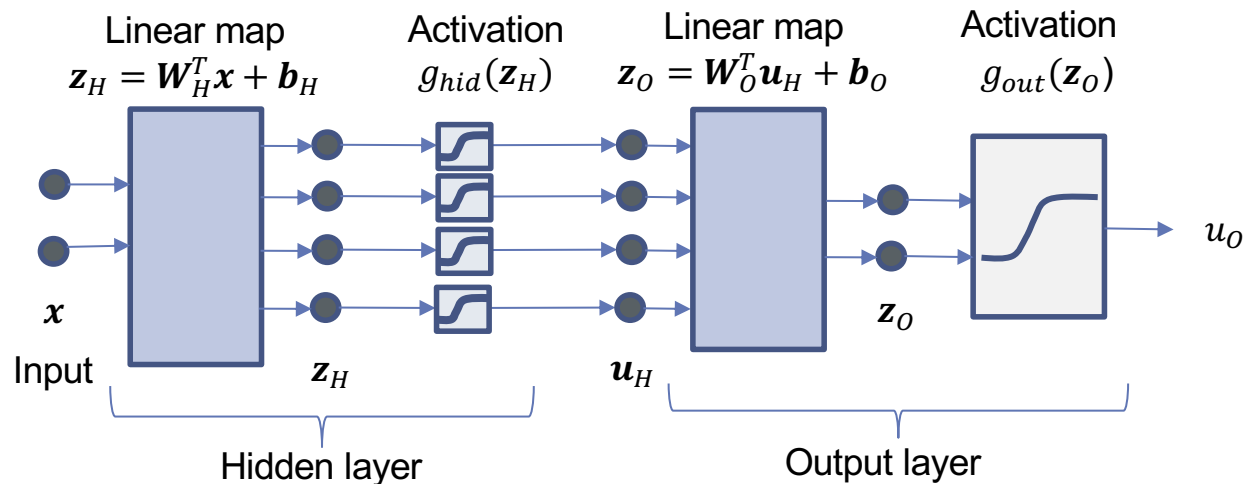
Ehsan Aryafar

earyafar@pdx.edu


<http://web.cecs.pdx.edu/~aryafare/ML.html>

Recall: General Neural Net Block Diagram

- Hidden layer: $\mathbf{z}_H = \mathbf{W}_H^T \mathbf{x} + \mathbf{b}_H$, $\mathbf{u}_H = g_{hid}(\mathbf{z}_H)$
- Output layer: $\mathbf{z}_O = \mathbf{W}_O^T \mathbf{u}_H + \mathbf{b}_O$, $u_O = g_{out}(\mathbf{z}_O)$




Recall: Selecting the Output Activation




Target	Num output units $=\dim(u_o) = \dim(z_o)$	Output activation $u_o = g_{out}(z_o)$	Interpretation
Binary classification	1	$u_o = \text{sigmoid}(z_o)$	$u_o = P(y = 1 x)$
K -class classification	K	$u_o = \text{softmax}(z_o)$	$u_{o,k} = P(y = k x)$
Regression with K outputs	K	$u_o = z_o$	$u_{o,k} = \hat{y}_k$

Recall: Selecting the Hidden Activation




- Two common choices
- **Sigmoid**:
 - $u_{H,k} = 1/(1 + \exp(-z_{H,k}))$
- **ReLU** (Rectified linear unit):
 - $u_{H,k} = \max\{0, z_{H,k}\}$


Sigmoid



ReLU



Recall: Stochastic Gradient Descent



- In each step:
 - Select random small “mini-batch”
 - Evaluate gradient on mini-batch
- For $t = 1$ to N_{steps}
 - Select random mini-batch $I \subset \{1, \dots, N\}$
 - Compute gradient approximation (only over mini-batch samples):
$$g^t = \frac{1}{|I|} \sum_{i \in I} \nabla L(x_i, y_i, \theta)$$
 - Update parameters:
$$\theta^{t+1} = \theta^t - \alpha^t g^t$$

Recall: Simple MNIST Neural Network

- 784 inputs, 100 hidden units, 10 outputs

```
nin = Xtr.shape[1] # dimension of input data
nh = 100          # number of hidden units
nout = int(np.max(ytr)+1) # number of outputs = 10 since there are 10 classes
model = Sequential()
model.add(Dense(units=nh, input_shape=(nin,), activation='sigmoid', name='hidden'))
model.add(Dense(units=nout, activation='softmax', name='output'))
```

```
model.summary()
```


Layer (type)	Output Shape	Param #
hidden (Dense)	(None, 100)	78500
output (Dense)	(None, 10)	1010

=====
 Total params: 79,510
 Trainable params: 79,510
 Non-trainable params: 0

Recall: Fitting the Model

- Run for 20 epochs, ADAM optimizer, batch size = 100
- Final accuracy = 0.972
- Not great, but much faster than SVM. Also, CNNs do better.

```
opt = optimizers.Adam(lr=0.001) # beta_1=0.9, beta_2=0.
model.compile(optimizer=opt,
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
```



```
model.fit(Xtr, ytr, epochs=10, batch_size=100, validation_data=(Xts,yts))
```

```
Epoch 7/10
50000/50000 [=====] - 3s - loss: 0.0474 - acc: 0.9868 - val_loss: 0.0886 - val_ac
c: 0.9717
Epoch 8/10
50000/50000 [=====] - 3s - loss: 0.0440 - acc: 0.9884 - val_loss: 0.0875 - val_ac
c: 0.9718
Epoch 9/10
50000/50000 [=====] - 2s - loss: 0.0393 - acc: 0.9903 - val_loss: 0.0872 - val_ac
c: 0.9732
Epoch 10/10
50000/50000 [=====] - 3s - loss: 0.0381 - acc: 0.9901 - val_loss: 0.0875 - val_ac
c: 0.9718
```

Recall: Initialization and Data Normalization

- Solution by gradient descent algorithm depends on the initial weights
- Typically, weights are set to random values near zero.
- Small weights make the network behave like linear classifier.
 - Hence model starts out nearly linearly
 - Becomes nonlinear as weights increase during the training process.
- Starting with large weights often lead to poor results.
- **Normalizing data to zero mean and unit variance**
 - **Allows all input dimensions be treated equally and facilitate better convergence.**
- With normalized data, it is typical to initialize the weights to be uniform in $[-0.7, 0.7]$ [ESL]

Recall: Regularization

- To avoid the weights get too large, can add a penalty term explicitly, with regularization level λ

- Ridge penalty

$$R(\theta) = \sum_{d,m} w_{H,d,m}^2 + \sum_{m,k} w_{O,m,k}^2 = \|w_H\|^2 + \|w_O\|^2$$

- Total loss

$$L_{reg}(\theta) = L(\theta) + \lambda R(\theta)$$

- Change in gradient calculation
- Typically used regularization
 - L2 = Ridge: Shrink the sizes of weights
 - L1: Prefer sparse set of weights
 - L1-L2: use a combination of both

Recall: Regularization in Keras

- `kernel_regularizer`: instance of `keras.regularizers.Regularizer`
- `bias_regularizer`: instance of `keras.regularizers.Regularizer`
- `activity_regularizer`: instance of `keras.regularizers.Regularizer`

Activity regularization tries to make the output at each layer small or sparse.

Example

```
from keras import regularizers
model.add(Dense(64, input_dim=64,
                kernel_regularizer=regularizers.l2(0.01),
                activity_regularizer=regularizers.l1(0.01)))
```

Available penalties

```
keras.regularizers.l1(0.)
keras.regularizers.l2(0.)
keras.regularizers.l1_l2(0.)
```

Learning Objectives


- Identify cases to use dimensionality reduction
- Mathematically describe principal components representations of data
- Compute principal components via SVDs
- Compute PC components in python
- Add PCA transforms as a pre-processing step to classification and regression
- Implement low-rank transforms for recommender systems

Outline

- Why dimensionality reduction?
- Principal components and directions of variance
- Approximation with PCs
- Computing PCs via the SVD
- Face example in python
- Training models from PCs
- Low rank approximations and recommender systems

High-Dimensional Data

- Many data sets have very high dimension
- Training can be difficult
 - Especially when number of samples is small
 - Classifier needs many parameters



EEG


Ex: 32 channels x 1 kHz x 10s




Face recognition with high-resolution images

Problems with High-Dimensions

- Consider face recognition
- Input is high-dimensional
 - Esp. for high resolution image
- Resulting classifier:
 - Requires many parameters
 - Difficult to train
 - Needs many samples
 - Computationally complex



Dimensionality Reduction




- Dimensionality reduction:
 - Reduce the input dimension to lower dimensional representation
- Can build simpler classifier
- Low-dimensional representation also good for:
 - Visualizing data
 - Clustering and other unsupervised tasks
 - Finding underlying structure of the data

Outline

- Dimensionality reduction
- Principal components and directions of variance
- Approximation with PCs
- Computing PCs via the SVD
- Face example in python
- Training models from PCs
- Low rank approximations and recommender systems

Data Definitions

- Given data: $\mathbf{x}_i, i = 1, \dots, N$
 - Each sample has p features: $\mathbf{x}_i = (x_{i1}, \dots, x_{ip})$
 - Represent as an $N \times p$ matrix
- Unsupervised learning
 - Samples do not have a label
 - Or we choose to ignore the label for now
- Dimension p is large
- How do we reduce the dimension?



Projections

- PCA reduces dimensionality by “projecting” data to a lower dim subspace
- **Projection:** Given vectors \mathbf{z} and \mathbf{v} , the projection of \mathbf{z} onto \mathbf{v} is:


$$\hat{\mathbf{z}} = \text{Proj}_{\mathbf{v}}(\mathbf{z}) = \alpha \mathbf{v}, \quad \alpha = \frac{\mathbf{v}^T \mathbf{z}}{\mathbf{v}^T \mathbf{v}} = \frac{\|\mathbf{z}\|}{\|\mathbf{v}\|} \cos \theta$$

- α = coefficient of the projection

- **Theorem:** $\text{Proj}_{\mathbf{v}}(\mathbf{z})$ is closest point in V to \mathbf{z} :


$$\hat{\mathbf{z}} = \arg \min_{\mathbf{w} \in V} \|\mathbf{z} - \mathbf{w}\|^2$$

- $V = \{\alpha \mathbf{v} | \alpha \in R\}$ = vectors on the line spanned by \mathbf{v}



Maximal Directional Variance

- Given data: $x_i, i = 1, \dots, N$ and direction v with $\|v\| = 1$
- Let $z_i = v^T x_i$ = coefficient of the projection of x_i onto v
- Sample mean and variance in direction v is :
 - Sample mean $\bar{z} = \frac{1}{N} \sum_{i=1}^N z_i$
 - Sample variance $s_z^2 = \frac{1}{N} \sum_{i=1}^N (z_i - \bar{z})^2$



Problem: Find the direction v that maximizes the variance s_z^2

- Why?
 - Captures the most variation of the data
 - Provides the best vector for dimensionality reduction

Sample Covariance Matrix

- Sample mean of the data: $\bar{\mathbf{x}} = \frac{1}{N} \sum_{i=1}^N \mathbf{x}_i$
- Sample covariance matrix: Matrix \mathbf{Q} with components:

$$Q_{k\ell} = \frac{1}{N} \sum_{i=1}^N (x_{ik} - \bar{x}_k)(x_{i\ell} - \bar{x}_\ell)$$

- Covariance between feature k and ℓ in the dataset
- Matrix is $p \times p$
- Sample covariance is given by

$$\mathbf{Q} = \frac{1}{N} \sum_{i=1}^N (\mathbf{x}_i - \bar{\mathbf{x}})(\mathbf{x}_i - \bar{\mathbf{x}})^T = \frac{1}{N} \tilde{\mathbf{X}}^T \tilde{\mathbf{X}}$$

- $\tilde{\mathbf{X}}$ = data matrix with sample mean removed (rows: $\tilde{\mathbf{x}}_i = \mathbf{x}_i - \bar{\mathbf{x}}$)
- Compute sample covariance via a matrix product

Sample Covariance and Directional Variance

- Let $z_i = \mathbf{v}^T \mathbf{x}_i$ = coefficient of the projection of \mathbf{x}_i onto \mathbf{v}
- Can compute the sample mean and variance of z_i from $\bar{\mathbf{x}}$ and \mathbf{Q}

- Sample mean of the coefficients:

$$\bar{z} = \frac{1}{N} \sum_{i=1}^N z_i = \frac{1}{N} \sum_{i=1}^N \mathbf{v}^T \mathbf{x}_i = \mathbf{v}^T \left[\frac{1}{N} \sum_{i=1}^N \mathbf{x}_i \right] = \mathbf{v}^T \bar{\mathbf{x}}$$

- Sample variance of the coefficients:

$$\begin{aligned} s_z^2 &= \frac{1}{N} \sum_{i=1}^N (z_i - \bar{z})^2 = \frac{1}{N} \sum_{i=1}^N (\mathbf{v}^T (\mathbf{x}_i - \bar{\mathbf{x}}))^2 \\ &= \frac{1}{N} \sum_{i=1}^N \mathbf{v}^T (\mathbf{x}_i - \bar{\mathbf{x}}) (\mathbf{x}_i - \bar{\mathbf{x}})^T \mathbf{v} \\ &= \mathbf{v}^T \mathbf{Q} \mathbf{v} \end{aligned}$$

Maximizing Directional Variance


- From previous slide: Directional variance $s_Z^2 = \frac{1}{N} \sum_{i=1}^N (z_i - \bar{z})^2 = \mathbf{v}^T \mathbf{Q} \mathbf{v}$
- Maximizing directional variance can be formulated as an optimization problem:

$$\max_{\mathbf{v}} \mathbf{v}^T \mathbf{Q} \mathbf{v} \quad \text{s.t.} \quad \|\mathbf{v}\| = 1$$

- Let $\mathbf{v}_1, \dots, \mathbf{v}_p$ be the **eigenvectors** of \mathbf{Q} : $\mathbf{Q} \mathbf{v}_j = \lambda_j \mathbf{v}_j$
- Sort eigenvalues in descending order: $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_p$
 - Can show that eigenvalues are real and non-negative
- **Theorem**: Any local maxima of the variance directional is an eigenvector
 - $\mathbf{v} = \mathbf{v}_j$ for some j and $\mathbf{v}^T \mathbf{Q} \mathbf{v} = \lambda_j$
 - Proof below

Visualizing Principal Components

- **Principal components:** The eigenvectors of \mathbf{Q} , $\mathbf{v}_1, \dots, \mathbf{v}_p$
 - Always normalized $\|\mathbf{v}_j\| = 1$
 - Sorted by eigenvalues $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_p$
 - Each vector is of dimension p
- **Key property:** Vectors are orthogonal
 - $\mathbf{v}_j^T \mathbf{v}_k = 0$ if $j \neq k$
- **Represents directions of decreasing variance:**
 - \mathbf{v}_1 : PC 1 = Direction of max variance
 - \mathbf{v}_2 : PC 2 = Direction of second most variance
 - \mathbf{v}_3 : PC 3 = Direction of third most variance
 - ...



Proof PCs = Eigenvectors of \mathbf{Q} (Advance Concept)

- PC constrained optimization problem:

$$\max_{\mathbf{v}} \mathbf{v}^T \mathbf{Q} \mathbf{v} \quad \text{s.t.} \quad \|\mathbf{v}\| = 1$$

- Define **Lagrangian**: $L(\mathbf{v}, \lambda) = \mathbf{v}^T \mathbf{Q} \mathbf{v} - \lambda [\|\mathbf{v}\|^2 - 1]$

- At any local maxima:

$$\frac{\partial L}{\partial \mathbf{v}} = 0 \Rightarrow \mathbf{Q} \mathbf{v} - \lambda \mathbf{v} = \mathbf{0}$$

- This shows that \mathbf{v} is an eigenvector of \mathbf{Q}

Outline


- Dimensionality reduction
- Principal components and directions of variance
- Approximation with PCs
- Computing PCs via the SVD
- Face example in python
- Training models from PCs
- Low rank approximations and recommender systems

Low-Dimensional Representations

- Given data $x_i, i = 1, \dots, N$. Each $x_i \in \mathbb{R}^p$
- **Problem:** Find **basis vectors** $v_j, j = 1, \dots, d$ such that:


$$x_i \approx \hat{x}_i = \bar{x} + \sum_{j=1}^d \alpha_{ij} v_j$$

- Sample mean + linear combination of basis vectors
- $\alpha_i = (\alpha_{i1}, \dots, \alpha_{id})$ is an approximate **coordinates** of x_i in basis (v_1, \dots, v_d)
- Dimensionality reduction:
 - If $d \ll p$ we have represented x_i with a smaller number of coefficients.



Orthonormal Sets and Basis

- **Definition:** A set of vectors $\mathbf{v}_1, \dots, \mathbf{v}_d$ are an **orthonormal set** if:
 - $\|\mathbf{v}_j\| = 1$ for all j (unit length)
 - $\mathbf{v}_j^T \mathbf{v}_k = 0$ if $j \neq k$ (perpendicular to one another)
- **Matrix form:** If $\mathbf{V} = [\mathbf{v}_1 \dots \mathbf{v}_d]$, then $\mathbf{V}^T \mathbf{V} = \mathbf{I}_d$
- If $d = p$ then $\mathbf{v}_1, \dots, \mathbf{v}_p$ is called an **orthonormal basis**
 - \mathbf{V} is an **orthonormal matrix**
- **Key property:** the PCs form an orthonormal basis



Coefficients in an Orthonormal Basis

- Suppose $\mathbf{v}_1, \dots, \mathbf{v}_p$ is an orthonormal basis
- Given a vector \mathbf{z} , can write

$$\mathbf{z} = \sum_{j=1}^p \alpha_j \mathbf{v}_j, \quad \alpha_j = \mathbf{v}_j^T \mathbf{z}$$

- Simple expression for computing coefficients in an orthonormal basis
- Matrix form:

$$\boldsymbol{\alpha} = \mathbf{V}^T \mathbf{z}, \quad \mathbf{z} = \mathbf{V} \boldsymbol{\alpha}$$

Approximating the Data Matrix

- Given data $\mathbf{x}_i, i = 1, \dots, N$
- Let $\mathbf{v}_1, \dots, \mathbf{v}_p$ be the PCs
- Find coefficient expansion of each data sample:

$$\mathbf{x}_i = \bar{\mathbf{x}} + \sum_{j=1}^p \alpha_{ij} \mathbf{v}_j, \quad \alpha_{ij} = \mathbf{v}_j^T (\mathbf{x}_i - \bar{\mathbf{x}})$$

- Approximation with d coefficients:


$$\hat{\mathbf{x}}_i = \bar{\mathbf{x}} + \sum_{j=1}^d \alpha_{ij} \mathbf{v}_j$$

Geometry of Approximations

- Approximation can be interpreted geometrically
- Let V be set of all linear combinations

$$\sum_{j=1}^d \alpha_j \mathbf{v}_j$$

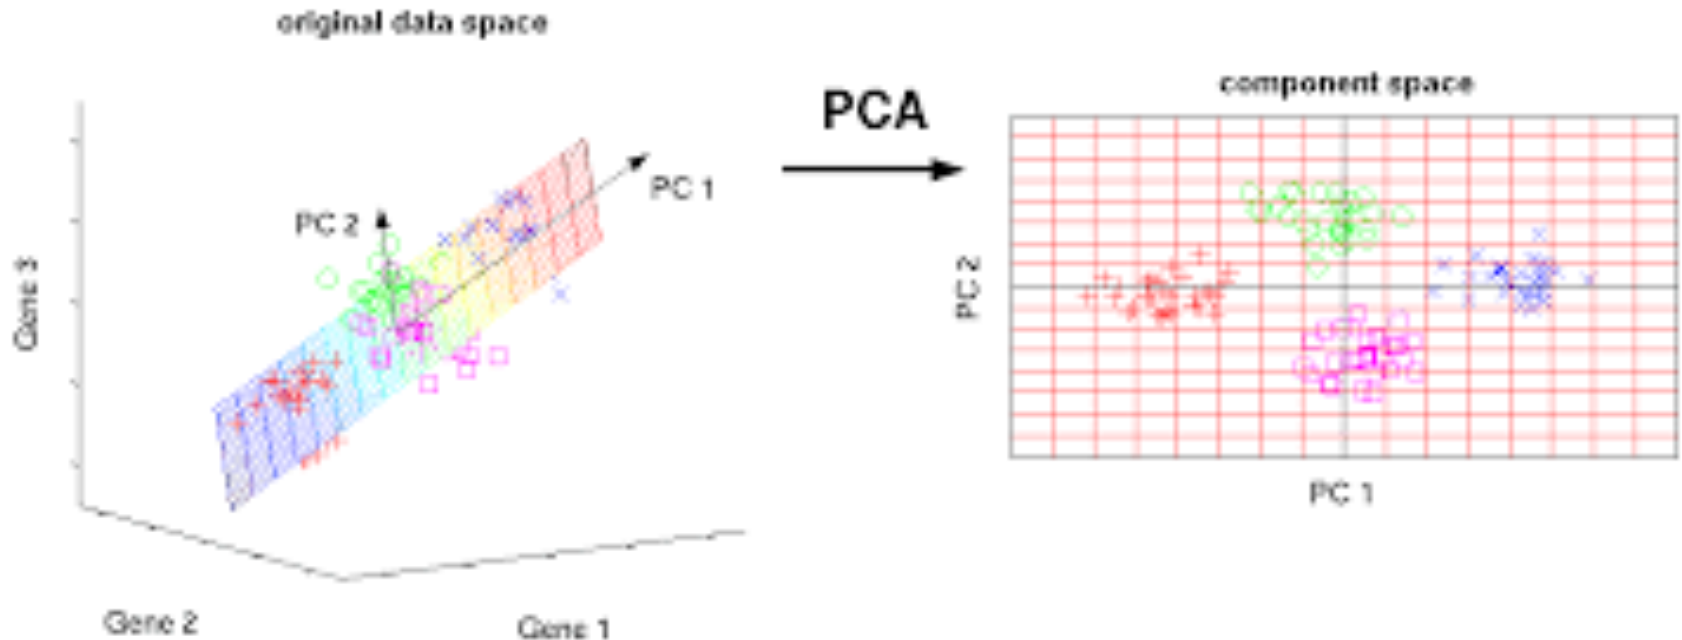
- V is a vector space
- Called the span of $\mathbf{v}_1, \dots, \mathbf{v}_d$
- $\hat{\mathbf{x}} - \bar{\mathbf{x}}$ is the closest vector in V to $\mathbf{x} - \bar{\mathbf{x}}$
 - Note the subtraction of the mean



Space spanned by $\mathbf{v}_1, \dots, \mathbf{v}_d$

Visualizing the Representation

- Finds a low-dimensional representation



Example Calculation

- Problem:

- Let $\mathbf{v}_1 = \frac{1}{\sqrt{2}}[1, 1, 0]$, $\mathbf{v}_2 = \frac{1}{\sqrt{6}}[1, -1, 2]$
- Show \mathbf{v}_1 and \mathbf{v}_2 are orthogonal

- Solution:

- $\mathbf{v}_1^T \mathbf{v}_1 = \frac{1}{2}(1^2 + 1^2 + 0^2) = 1$
- $\mathbf{v}_2^T \mathbf{v}_2 = \frac{1}{6}(1^2 + (-1)^2 + 2^2) = 1$
- $\mathbf{v}_1^T \mathbf{v}_2 = \frac{1}{\sqrt{2(3)}}(1(1) + 1(-1) + 0(2)) = 0$

Example Calculation Continued

- Problem:

- Let $\mathbf{v}_1 = \frac{1}{\sqrt{2}} [1, 1, 0]$, $\mathbf{v}_2 = \frac{1}{\sqrt{6}} [1, -1, 2]$ be two PCs
- Let $\bar{\mathbf{x}} = [0, 1, 2]$ be the mean of the data
- Find the approximation of data record $\mathbf{x} = [2, 4, 4]$ with the two PCs

- Solution:

- Subtract mean: $\mathbf{x} - \bar{\mathbf{x}} = [2, 3, 2]$
- Coeff on PC1: $\alpha_1 = \mathbf{v}_1^T (\mathbf{x} - \bar{\mathbf{x}}) = \frac{1}{\sqrt{2}} [2 + 3 + 0] = \frac{5}{\sqrt{2}}$
- Coeff on PC2: $\alpha_2 = \mathbf{v}_2^T (\mathbf{x} - \bar{\mathbf{x}}) = \frac{1}{\sqrt{6}} [2 - 3 + 4] = \frac{3}{\sqrt{6}}$
- Approximation: $\hat{\mathbf{x}} = \bar{\mathbf{x}} + \sum_{j=1}^d \alpha_j \mathbf{v}_j = [0, 1, 2] + \frac{5}{2} [1, 1, 0] + \frac{3}{6} [1, -1, 2] \approx [3, 3, 3]$

Average Approximation Error

- Let $\hat{\mathbf{x}}_i$ = approximation with d PCs
- Error in sample i :

$$\mathbf{x}_i - \hat{\mathbf{x}}_i = \sum_{j=d+1}^p \alpha_{ij} \mathbf{v}_j$$

- **Theorem:** Average error with a d PC approximation is:

$$\frac{1}{N} \sum_{i=1}^N \|\mathbf{x}_i - \hat{\mathbf{x}}_i\|^2 = \sum_{j=d+1}^p \lambda_j$$

- Sum of the smallest $p - d$ eigenvalues

Proportion of Variance (PoV)

- Total variance of data set:

$$\frac{1}{N} \sum_{i=1}^N \|\mathbf{x}_i - \bar{\mathbf{x}}\|^2 = \sum_{j=1}^p \lambda_j$$

- Average approximation error:

$$\frac{1}{N} \sum_{i=1}^N \|\mathbf{x}_i - \hat{\mathbf{x}}_i\|^2 = \sum_{j=d+1}^p \lambda_j$$

- The **proportion of variance** explained by d PCs is:

$$PoV(d) = \frac{\sum_{j=1}^d \lambda_j}{\sum_{j=1}^p \lambda_j}$$

- Measure of approximation error in using d PCs


Example: suppose dataset with four dimensions

PC index	λ_i	POV(i)
1	10	10/14.3 \approx 0.70
2	4	14/14.3 \approx 0.98
3	0.2	14.2/14.3 \approx 0.99
4	0.1	14.3/14.3 = 1

Latent Representations

- Each record is of the form: $x_i \approx \bar{x} + \sum_{j=1}^d \alpha_{ij} v_j$
- Variance in x_i explained by small number of “latent components”
 - Coefficients α_{ij} are the latent representations of x_i
- Example:
 - x_i = list of movie preferences for customer i
 - Movie preferences are highly correlated.
 - Could be explained by small number of components (action, romance, presence of stars, ...)
 - PCA can be used to find these out

Example: USArrests




- Arrests per capita in four categories
 - One record per US state
- Visualize PCA in a biplot
 - See the scores (i.e. coefficients of each state)
 - Overlay loading plot (PC vectors)
- Fig from ESL 10.1

Outline

- Dimensionality reduction
- Principal components and directions of variance
- Approximation with PCs
- Computing PCs via the SVD
- Face example in python
- Training models from PCs
- Low rank approximations and recommender systems

Singular Value Decomposition

- **SVD**: Powerful method in linear algebra
- Given a matrix A :
 - Decomposes the matrix into a product: $A = U\Sigma V^T$
 - Provides orthonormal bases of the input and output spaces
 - Multiplication of A is equivalent to scaling in that basis
- For PCA:
 - Identifies low rank subspaces for data
 - Computes coefficients in that subspace



Singular Value Decomposition Defined


- Given matrix $A \in \mathbb{F}^{n \times d}$ (F is for field, just consider R)
 - For PCA, this will be a scaled version of the data matrix
- SVD is $A = U\Sigma V^T$, where
 - $U \in \mathbb{F}^{n \times r}$, columns are orthonormal
 - $V \in \mathbb{F}^{d \times r}$, columns are orthonormal
 - $\Sigma = \text{diag}(s_1, \dots, s_r)$, sorted $s_1 \geq s_2 \geq \dots \geq s_r \geq 0$.
 - Called the singular values
- All matrices have an SVD
 - Matrices do not have to be square.
- Number of singular values $r \leq \min(n, d)$

Economy vs. Full SVD

- Suppose $A \in \mathbb{R}^{n \times d}$ with rank $r \leq \min\{n, d\}$
- Two types of SVDs
- **Economy SVD**: $A = USV^*$
 - $U \in \mathbb{F}^{n \times r}$, columns are orthonormal
 - $V \in \mathbb{F}^{d \times r}$, columns are orthonormal
 - $\Sigma \in \mathbb{F}^{r \times r}$ diagonal $\Sigma = \text{diag}(s_1, \dots, s_r)$,
- **Full SVD**: $A = USV^*$
 - $U \in \mathbb{F}^{n \times n}$, columns are an orthonormal basis of \mathbb{R}^n
 - $V \in \mathbb{F}^{d \times d}$, columns are an orthonormal basis of \mathbb{R}^d
 - $\Sigma \in \mathbb{F}^{n \times d}$ with diagonal upper left $\Sigma = \begin{bmatrix} \hat{\Sigma} & 0 \\ 0 & 0 \end{bmatrix}$

SVD Visualized

- Pink matrices represent “economy” SVD
- Blue represent “full SVD”



Example

- Let $A = \begin{bmatrix} 1 & 0 & 0 & 0 & 2 \\ 0 & 0 & 3 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 2 & 0 & 0 & 0 \end{bmatrix}$


- Then can check that $A = U\Sigma V^*$

$$U = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & -1 \\ 1 & 0 & 0 & 0 \end{bmatrix} \quad \Sigma = \begin{bmatrix} 2 & 0 & 0 & 0 & 0 \\ 0 & 3 & 0 & 0 & 0 \\ 0 & 0 & \sqrt{5} & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix} \quad V^* = \begin{bmatrix} 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ \sqrt{0.2} & 0 & 0 & 0 & \sqrt{0.8} \\ 0 & 0 & 0 & 1 & 0 \\ -\sqrt{0.8} & 0 & 0 & 0 & \sqrt{0.2} \end{bmatrix}$$

- Also verify that $UU^* = I_5$ and $VV^* = I_5$
- But, in general, use a computer to compute SVD

Geometric Interpretation (Advanced)

- Let $A = U\Sigma V^*$ and $y = Ax$
- Consider a transformed space
 - $\mathbf{w} = V^* \mathbf{x} = [w_1, \dots, w_N]$ coefficients in input basis $V = [v_1, \dots, v_N]$
 - $\mathbf{z} = U^* \mathbf{y} = [z_1, \dots, z_M]$: coefficients in output basis $U = [u_1, \dots, u_M]$
- Then: $\mathbf{z} = \Sigma \mathbf{w}$ so $z_i = \sigma_i w_i$
- Each input direction \mathbf{v}_i is mapped to $\sigma_i \mathbf{u}_i$
- Consequence:
 - SVD finds orthonormal bases U, V such that matrix A is a linear scaling in each basis vector



Example Problem

- Suppose that $A = U\Sigma V^* \in \mathbb{R}^{3 \times 4}$ with $\Sigma = \text{diag}(3, 0.2, 0, 0)$
- If $\mathbf{x} = 2\mathbf{v}_1 + 3\mathbf{v}_2 + 4\mathbf{v}_3 + 5\mathbf{v}_4$ find $\mathbf{y} = A\mathbf{x}$ in terms of basis $\mathbf{u}_1, \mathbf{u}_2, \mathbf{u}_3$
- Solution:
 - $A\mathbf{v}_i = \sigma_i \mathbf{u}_i$ for all i
 - Therefore,

$$\begin{aligned}\mathbf{y} = A\mathbf{x} &= 2A\mathbf{v}_1 + 3A\mathbf{v}_2 + 4A\mathbf{v}_3 + 5A\mathbf{v}_4 \\ &= 2(3)\mathbf{u}_1 + 3(0.2)\mathbf{u}_2 + 4(0)\mathbf{u}_3 \\ &= 6\mathbf{u}_1 + 0.6\mathbf{u}_2\end{aligned}$$

Computing the SVD in Python

- Random matrix

```
# Create some random matrix
A = np.random.normal(0,1,(100,10))
```

- Full SVD

```
# Full SVD
U,s,Vtr = np.linalg.svd(A)
```

```
A.shape = (100, 10)
U.shape = (100, 100)
s.shape = (10,)
Vtr.shape = (10, 10)
```

- Economy SVD

```
# Economy SVD
U,s,Vtr = np.linalg.svd(A, full_matrices=False)
```

```
U.shape = (100, 10)
s.shape = (10,)
Vtr.shape = (10, 10)
```

- Reconstruction:

```
# Recovers back A
Ahat = (U*s[None,:]).dot(Vtr)
```


Computing the PCA via SVD

- Let $A = \frac{1}{\sqrt{N}} \tilde{X}$ = scaled data matrix with sample mean removed.
- Take SVD: $A = U\Sigma V^T$
- Properties:
 - Sample covariance matrix is $Q = \frac{1}{N} \tilde{X}^T \tilde{X} = A^T A = V\Sigma U^T U\Sigma V^T = V\Sigma^2 V^T$
 - Eigenvalues of Q = squared singular values of A
 - PCs are v_j , columns of V
 - Coefficients are $Z = \tilde{X}V = \sqrt{N}AV = \sqrt{N}U\Sigma$
- Hence, SVD provides PCs, eigenvalues coefficients, Z in the PCA representation.

Outline

- Dimensionality reduction
- Principal components and directions of variance
- Approximation with PCs
- Computing PCs via the SVD
- Face recognition using PCA in python
- Training models from PCs
- Low rank approximations and recommender systems

Example: Face Recognition



- Face recognition challenges:
 - Face images can be high-dimensional
 - We will use $50 \times 37 = 1850$ pixels
- Applying PCA:
 - Should be few degrees of freedom
 - Can transform to lower dimensional representations
- Data Labelled Faces in the Wild project
 - <http://vis-www.cs.umass.edu/lfw>
 - Large collection of faces (13000 images)
 - Taken from web articles about 20 years ago

Labeled Faces in the Wild Home



Loading the Data

- Built-in routines to load data from sk-learn package
- Can take several minutes the first time (Be patient)

```
from sklearn.datasets import fetch_lfw_people  
lfw_people = fetch_lfw_people(min_faces_per_person=70, resize=0.4)
```

```
2016-11-14 14:15:30,862 Downloading LFW metadata: http://vis-www.cs.umass.edu/lfw/pairsDevTrain.txt  
2016-11-14 14:15:30,958 Downloading LFW metadata: http://vis-www.cs.umass.edu/lfw/pairsDevTest.txt  
2016-11-14 14:15:31,028 Downloading LFW metadata: http://vis-www.cs.umass.edu/lfw/pairs.txt  
2016-11-14 14:15:31,294 Downloading LFW data (~200MB): http://vis-www.cs.umass.edu/lfw/lfw-funneled.tgz
```

```
Image size      = 50 x 37 = 1850 pixels  
Number faces    = 1288  
Number classes  = 7
```

Plotting the Data

- Some example faces
- You may be too young to remember them all

Colin Powell



Colin Powell



Hugo Chavez



George W Bush



```
def plt_face(x):  
    h = 50  
    w = 37  
    plt.imshow(x.reshape((h, w)), cmap=plt.cm.gray)  
    plt.xticks([])  
    plt.yticks([])  
  
I = np.random.permutation(n_samples)  
plt.figure(figsize=(10,20))  
nplt = 4;  
for i in range(nplt):  
    ind = I[i]  
    plt.subplot(1,nplt,i+1)  
    plt_face(X[ind])  
    plt.title(target_names[y[ind]])
```

Computing the PCA


```
npix = h*w  
Xmean = np.mean(X,0)  
Xs = X - Xmean[None,:]
```

```
U,S,Vtr = np.linalg.svd(Xs, full_matrices=False)
```

```
from sklearn.decomposition import PCA  
  
# Construct the PCA object  
pca = PCA(n_components=ncomp, svd_solver='randomized', whiten=True)  
  
# Fit the PCA components on the entire dataset  
pca.fit(X)
```

- Manually compute the PCs with SVD
 - Remove the mean
 - Use broadcasting
 - Compute the SVD
- Use sklearn builtin PCA function
 - Construct a PCA object
 - Call fit: Computes mean and PC components
 - Stores values internally in the pca class

Finding the PoV



- Most variance explained in about 400 components
- Some reduction

```
lam = S**2
PoV = np.cumsum(lam)/np.sum(lam)

plt.plot(PoV)
plt.grid()
plt.axis([1,n_samples,0, 1.1])
plt.xlabel('Number of PCs', fontsize=16)
plt.ylabel('PoV', fontsize=16)
```

Plotting Approximations

```
nplt = 2          # number of faces to plot
ds = [0,5,10,20,100] # number of SVD approximations
use_pca = True    # True=Use sklearn reconstruction, else use SVD
```


```
# Loop over figures
iplt = 0
for ind in inds:
    for d in ds:
        plt.subplot(nplt,nd+1,iplt+1)
        if use_pca:
            # Zero out coefficients after d.
            # Note, we need to copy to not overwrite the coefficients
            Zd = np.copy(Z[ind,:])
            Zd[d:] = 0
            Xhati = pca.inverse_transform(Zd)
        else:
            # Reconstruct with SVD
            Xhati = (U[ind,:d]*S[None,:d]).dot(Vtr[:,d,:]) + Xmean

        plt_face(Xhati)
        plt.title('d={0:d}'.format(d))
        iplt += 1

# Plot the true face
plt.subplot(nplt,nd+1,iplt+1)
plt_face(X[ind,:])
plt.title('Full')
iplt += 1
```


- Reconstruction using sklearn method
 - Uses the `inverse_transform` method to get back values
- Reconstruction using SVD
 - Note use of broadcasting

Plotting the Approximations



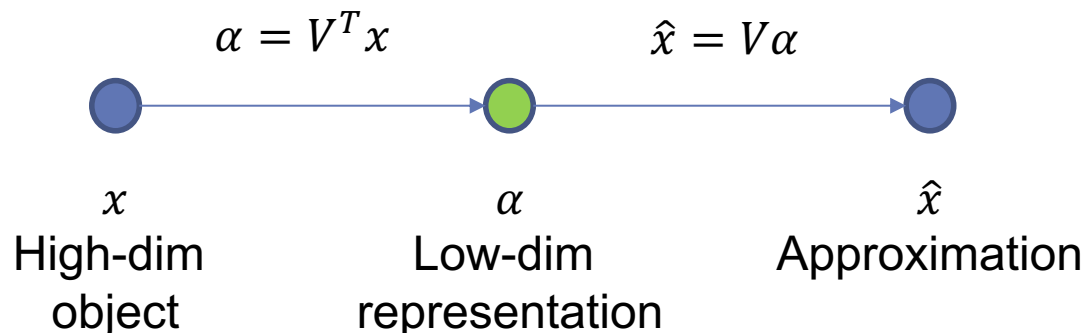
Plotting the PCs

- The PCs can be plotted as well




State-of-the-Art: Auto-Encoders

- PCA is a simple example of an **autoencoder**
- Tries to find low-dim representation
- Restricted to linear transforms
- Not very good for images and complex data



Deep Auto-Encoders


- Can use deep networks for learning complex latent representations and their inverses
 - http://www.cc.gatech.edu/~hays/7476/projects/Avery_Wenchen/
 - <https://swarbrickjones.wordpress.com/2016/01/13/enhancing-images-using-deep-convolutional-generative-adversarial-networks-dcgans/> (Code in Theano not tensorflow)



Outline


- Dimensionality reduction
- Principal components and directions of variance
- Approximation with PCs
- Computing PCs via the SVD
- Face example in python
- Training models from PCs
- Low rank approximations and recommender systems

Classification Using PCs



- Many problems: Dimensionality of data x is too large
 - Classifier in original space will have too many parameters
- Key idea:
 - Learn a dimension reducing transform via PCA: $z = f(x)$
 - Train classifier on low-dim transform $\hat{y} = g(z)$

Why This Would Work?



- PCA works if:
classes are separable
in transformed domain
- Example to left:
 - MNIST digits plotted in
two PCs
 - Can mostly separate the
classes

Training and Testing

- Split data in training and test: $X_{tr}, y_{tr}, X_{ts}, y_{ts}$
- Fit PCA transform on $Z = g(X)$ on training data X_{tr}
 - Do not include test data in PCA fit!
 - Many students make this mistake
- Transform training and test:
 - $Z_{tr} = g(X_{tr}), Z_{ts} = g(X_{ts})$
- Fit classifier $\hat{y} = f(z)$ on transformed training data (Z_{tr}, y_{tr})
- Predict classifier on transformed test data: $\hat{y}_{ts} = f(Z_{ts})$
- Score error rate / MSE on test data: $\epsilon = \frac{1}{N} \#\{\hat{y}_{ts}^i \neq y_{ts}^i\}$

How low of a dimension should I choose?

Cross-Validation

- To find number of PCs and other parameters use cross-validation (you can also use k-fold validation)
- Split data in training and test: $X_{tr}, y_{tr}, X_{ts}, y_{ts}$
- For each set of parameters:
 - Fit PCA transform on $Z = g(X, \text{numPCs})$ on training data X_{tr}
 - Transform training and test: $Z_{tr} = g(X_{tr})$, $Z_{ts} = g(X_{ts})$
 - Fit classifier $\hat{y} = f(z)$ on transformed training data (Z_{tr}, y_{tr})
 - Predict classifier on transformed test data: $\hat{y}_{ts} = f(Z_{ts})$
 - Score (e.g. error rate / MSE) on test data: $\epsilon = \frac{1}{N} \#\{\hat{y}_{ts}^i \neq y_{ts}^i\}$
- Select the parameters with lowest score
 - Number of PCs to use
 - Classifier may have some parameters too, e.g., gamma in RBF for SVM

Example: SVM classification with PCAs

```

npc_test = [25,50,75,100,200]
gam_test = [1e-3,4e-3,1e-2,1e-1]
C = 100
n0 = len(npc_test)
n1 = len(gam_test)
acc = np.zeros((n0,n1))
acc_max = 0

for i0, npc in enumerate(npc_test):

    # Fit PCA on the training data
    pca = PCA(n_components=npc, svd_solver='randomized', whiten=True)
    pca.fit(Xtr)

    # Transform the training and test
    Ztr = pca.transform(Xtr)
    Zts = pca.transform(Xts)

    for i1, gam in enumerate(gam_test):

        # Fitting on the transformed training data
        svc = SVC(C=c, kernel='rbf', gamma = gam)
        svc.fit(Ztr, ytr)

        # Predict on the test data
        yhat = svc.predict(Zts)


        # Compute the accuracy
        acc[i0,i1] = np.mean(yhat == yts)
        print('npc=%d gam=%12.4e acc=%12.4e' % (npc,gam,acc[i0,i1]))

```

- Parameters to search
 - Number of PCs and gamma for RBF in SVM
 - C param in SVM fixed to 100
- Fit on the training data.
 - This is in the loop!
- Transform the data
- Fit classifier on transformed training data
- Test on the transformed test data
- Score on test data

Example: Parameter Search

- Search over:
 - Number of PCs $\in \{25, 50, 75, 100, 200\}$
 - $\gamma \in \{0.001, 0.004, 0.01, 0.1\}$
- Plotted is the test accuracy
- Best test accuracy $\approx 85\%$
- Original data has 1850 dimension, but 75 PCs is optimal! Large reduction!
- More PCs reduces accuracy



Optimal num PCs = 75
Optimal gamma = 0.010000

Examples

- Correct images

George W Bush George W Bush



George W Bush George W Bush



Original Reduced

- Error images

Tony Blair George W Bush



Gerhard Schroeder George W Bush



Original Reduced

Outline

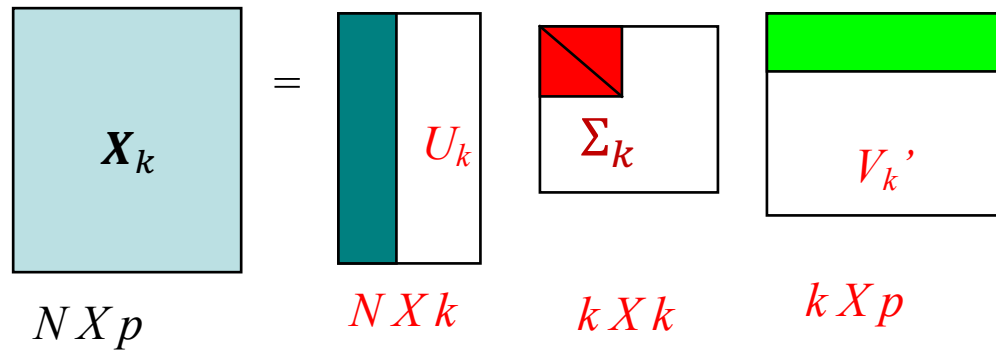
- Dimensionality reduction
- Principal components and directions of variance
- Approximation with PCs
- Computing PCs via the SVD
- Face example in python
- Training models from PCs
- Low rank approximations and recommender systems

Low-Rank Approximations

- SVD can be used for a low-rank approximation
- SVD can be written: $\mathbf{X} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T = \sum_{j=1}^r \alpha_j \mathbf{u}_j \mathbf{v}_j^T$
- Consider k –term approximation: $\mathbf{X}_k = \sum_{j=1}^k \alpha_j \mathbf{u}_j \mathbf{v}_j^T$
- Properties:
 - \mathbf{X}_k is rank k
 - $\mathbf{X}_k = \mathbf{U}_k \mathbf{\Sigma}_k \mathbf{V}_k^T$
 - Error is $\|\mathbf{X} - \mathbf{X}_k\|_F^2 = \sum_i \sum_j (X_{ij} - X_{k,ij})^2 = \sum_{j=k+1}^r \alpha_j^2$
 - If s_{k+1}, \dots, s_r is small, then matrix is well approximated by rank k matrix

The next slides just give some high level ideas!

Low-Rank Approximation Visualized



The diagram illustrates the Low-Rank Approximation $X_k = U_k \Sigma_k V_k'$. It shows four matrices represented by colored rectangles:

- A light blue rectangle labeled X_k with dimensions $N \times p$ below it.
- An equals sign.
- A teal rectangle labeled U_k with dimensions $N \times k$ below it.
- A white rectangle with a red diagonal line and a red square in the top-left corner, labeled Σ_k with dimensions $k \times k$ below it.
- A white rectangle with a green horizontal bar at the top, labeled V_k' with dimensions $k \times p$ below it.

- Can show: Reconstructed matrix X_k is optimal rank k approximation

Recommender Systems

- How do you recommend a movie to a user?
- MovieLens dataset:
 - Get past ratings from users
 - Make recommendations for future

t[3]:

	movielid	title	genres
0	1	Toy Story (1995)	Adventure Animation Children Comedy Fantasy
1	2	Jumanji (1995)	Adventure Children Fantasy
2	3	Grumpier Old Men (1995)	Comedy Romance
3	4	Waiting to Exhale (1995)	Comedy Drama Romance
4	5	Father of the Bride Part II (1995)	Comedy

Bo Burnham: Words, Words, Words (TV Special 2010)

TV MA Documentary | Comedy | Music


★★★★★ 8.2/10

The internet (and soon to be movie, TV, radio, etc.) phenomenon, Bo Burnham, brings you his first one-hour stand-up special "Bo Burnham: Words, Words, Words" from the House of Blues in Boston.

Director: Shannon Hartman
Stars: Bo Burnham

Ratings Matrix

- Data can be represented as **ratings matrix**
 - Users x movies
- **Problem**: Most users have only rated a small fraction
- Need to estimate unseen entries
 - Very sparse
- How can we complete this matrix



Name	Dates	Users	Movies	Ratings	Density
ML Latest	'95 – '16	247,753	34,208	22,884,377	0.003%
ML Latest Small	'96 – '16	668	10,329	105,339	0.015%

Latent Factor Model for Ratings

- **Idea:** Ratings for movies dependent on small number of **latent** factors
 - E.g. Action, famous actors, genre, ...
- Mathematically model as:

$$R_{ij} \approx \hat{R}_{ij} = b_i^u + b_j^m + \sum_{k=1}^K A_{ik} B_{jk}$$

- R_{ij} = Rating of movie j by user i
- b_i^u = Bias of user i
- b_j^m = Bias of movie j
- K = number of latent factors. Typically small $K \ll N_{user}, N_{movies}$
- A_{ik} = Preference of user i to factor k
- B_{jk} = Component of factor k in movie j